

Operating System

Tutorial 10

Deadlock Avoidance (Dijkstra/Banker's Algorithm)

Deadlock

- A **set** of blocked processes **each holding** some resources and **waiting to acquire** some resources **held by another process in the set**.
- Example (hardware resources):
 - System has 2 tape drives.
 - P_1 and P_2 each hold one tape drive and each needs another one.
- Example (“software” resource):
 - semaphores A and B , initialized to 1

| P_0 | P_1 |
|-------------------|------------------|
| <i>wait (A);</i> | <i>wait(B)</i> |
| <i>wait (B);</i> | <i>wait(A)</i> |
| <i>signal(B);</i> | <i>signal(A)</i> |
| <i>signal(A);</i> | <i>signal(B)</i> |

The (4) Conditions for Deadlock

There are three **policy** conditions that must hold **for a deadlock to be possible** (the “necessary conditions”):

1. **Mutual Exclusion**

- only one process at a time may use a resource

2. **Hold-and-Wait**

- a process may **hold** allocated resources while **awaiting** assignment of others


3. **No Preemption**

- a resource can be **released only voluntarily** by the process holding it, after that process has completed its task

- ..and a fourth condition **which must actually arise** to make a deadlock happen

The Conditions for Deadlock


- In the presence of these necessary conditions, one more condition must arise for deadlock to actually occur:
- 4. Circular Wait
 - a **closed chain** of processes exists, such that each process **holds at least one** resource needed by the **next process in the chain**



The Conditions for Deadlock

- Deadlock occurs **if and only if** the **circular wait** condition is **unresolvable**
- The circular wait condition is unresolvable if the first **3 policy** conditions hold
- So the 4 conditions taken together constitute necessary and sufficient conditions for deadlock

Resource Request Allocation Graph:



Multiple instances of resources R_1 and R_2 : P_2 and P_4 can complete, freeing up resources for P_1 and P_3

Approaches to Handling Deadlocks

- Deadlock **Prevention**
 - disallow 1 of the 3 necessary conditions of deadlock occurrence, or prevent the circular wait condition from happening
- Deadlock **Avoidance**
 - do not grant a resource request if this allocation **might** lead to deadlock
- Deadlock **Detection and Recovery**
 - grant resource requests when possible, but periodically **check** for the presence of deadlock and then take action to **recover** from it

Deadlock Avoidance

- Two approaches:
 - do not **start** a process if its total demand might lead to deadlock: (“**Process Initiation Denial**”), or
 - do not grant an **incremental** resource request if this allocation could lead to deadlock: (“**Resource Allocation Denial**”)
- In both cases: **maximum** requirements of each resource must be stated in advance

Resource Allocation Denial

- A Better Approach:
 - Grant **incremental resource requests** if we can prove that this leaves the system in a state in which **deadlock cannot occur**.
 - Based on the concept of a “safe state”.
- Banker’s Algorithm: (developed by Edsger Dijkstra)
 - **Tentatively** grant each resource request
 - Analyze resulting system state to see if it is “**safe**”.
 - If **safe**, grant the request
 - if **unsafe** refuse the request (undo the tentative grant)
 - block the requesting process until it is safe to grant it.

Data Structures for the Banker's Algorithm

Let n = number of processes,

m = number of resource types

- **Available**: Vector of length m . If $\text{Available}[j] = k$, there are k instances of resource type R_j **currently available**
- **Max**: $n \times m$ matrix. If $\text{Max}[i,j] = k$, then process P_i will request **at most k** instances of resource type R_j .
- **Alloc**: $n \times m$ matrix. If $\text{Alloc}[i,j] = k$ then P_i is currently allocated (i.e. holding) k instances of R_j .
- **Need**: $n \times m$ matrix. If $\text{Need}[i,j] = k$, then P_i may need **k more** instances of R_j to complete its task.

$$\text{Need}[i,j] = \text{Max}[i,j] - \text{Alloc}[i,j].$$

Safety Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively. Initialize:
Work := *Available*
Finish [i] == *false* for $i = 1, 2, \dots, n$.
2. Find an i such that both:
Finish [i] == *false*
 $Need_i \leq Work$
If no such i exists, go to step 4.
3. *Work* := *Work* + *Allocation* _{i}
(Resources freed when process completes!)
Finish [i] := *true*
go to step 2.
4. If *Finish* [i] = *true* for all i , then the system is in a safe state.

Resource-Request Algorithm for Process P_i

Request_i = request vector for P_i .

Request_i [j] = k means process P_i wants k instances of resource type R_j .

1. If **Request_i ≤ Need_i** go to step 2. Otherwise, **error** (**process exceeded its maximum claim**).

2. If **Request_i ≤ Available**, go to step 3. Otherwise P_i must wait, (**resources not available**).

3. “Allocate” requested resources to P_i as follows:

Available := Available - Request_i

Alloc_i := Alloc_i + Request_i

Need_i := Need_i - Request_i

If **safe** \Rightarrow the resources are **allocated** to P_i .

If **unsafe** \Rightarrow **restore** the **old resource-**

Example of Banker's Algorithm

5 processes P_0 through P_4

3 resource types A (10 units), B (5 units), and C (7 units).

Snapshot at time T_0 :

| | <u>Allocation</u> | | | <u>Max</u> | | | <u>Available</u> | | |
|-------|-------------------|---|---|------------|---|---|------------------|---|---|
| | A | B | C | A | B | C | A | B | C |
| P_0 | 0 | 1 | 0 | 7 | 5 | 3 | 3 | 3 | 2 |
| P_1 | 2 | 0 | 0 | 3 | 2 | 2 | | | |
| P_2 | 3 | 0 | 2 | 9 | 0 | 2 | | | |
| P_3 | 2 | 1 | 1 | 2 | 2 | 2 | | | |
| P_4 | 0 | 0 | 2 | 4 | 3 | 3 | | | |

Example (cont)

Need = Max – Allocation

Need

A B C

P_0 7 4 3

P_1 1 2 2

P_2 6 0 0

P_3 0 1 1

P_4 4 3 1

The system is in a safe state since the sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfies safety criteria.

Now P_1 requests (1,0,2)

Check that Request \leq Available
(that is, $(1,0,2) \leq (3,3,2) \Rightarrow \text{true}$.)

| | <u>Allocation</u> | | | <u>Need</u> | | | <u>Available</u> | | |
|-------|-------------------|---|---|-------------|---|---|------------------|---|---|
| | A | B | C | A | B | C | A | B | C |
| P_0 | 0 | 1 | 0 | 7 | 4 | 3 | 2 | 3 | 0 |
| P_1 | 3 | 0 | 2 | 0 | 2 | 0 | | | |
| P_2 | 3 | 0 | 2 | 6 | 0 | 0 | | | |
| P_3 | 2 | 1 | 1 | 0 | 1 | 1 | | | |
| P_4 | 0 | 0 | 2 | 4 | 3 | 1 | | | |