



# Memory Management


CS303

Operating System

# Big Picture

The title 'Big Picture' is positioned on the left side of the slide. To its right, there are six circles arranged in a horizontal row. The first circle is solid light purple. The second circle is white with a light purple outline. The third circle is solid light purple. The fourth circle is white with a light purple outline. The fifth circle is solid light purple. The sixth circle is solid light purple.

- Up till now, we've focused on how multiple programs share the CPU
- Now: how do multiple programs share memory?





# Goals of Memory Management

- Allocate scarce memory resources among competing processes
  - While maximizing memory utilization and system throughput
- Provide a convenient abstraction for programming (and for compilers, etc.)
  - Hide sharing
  - Hide scarcity
- Provide isolation between processes


# Problem: Memory Relocation

- A program “sees” different memory addresses, depending on who else is running

3 MB

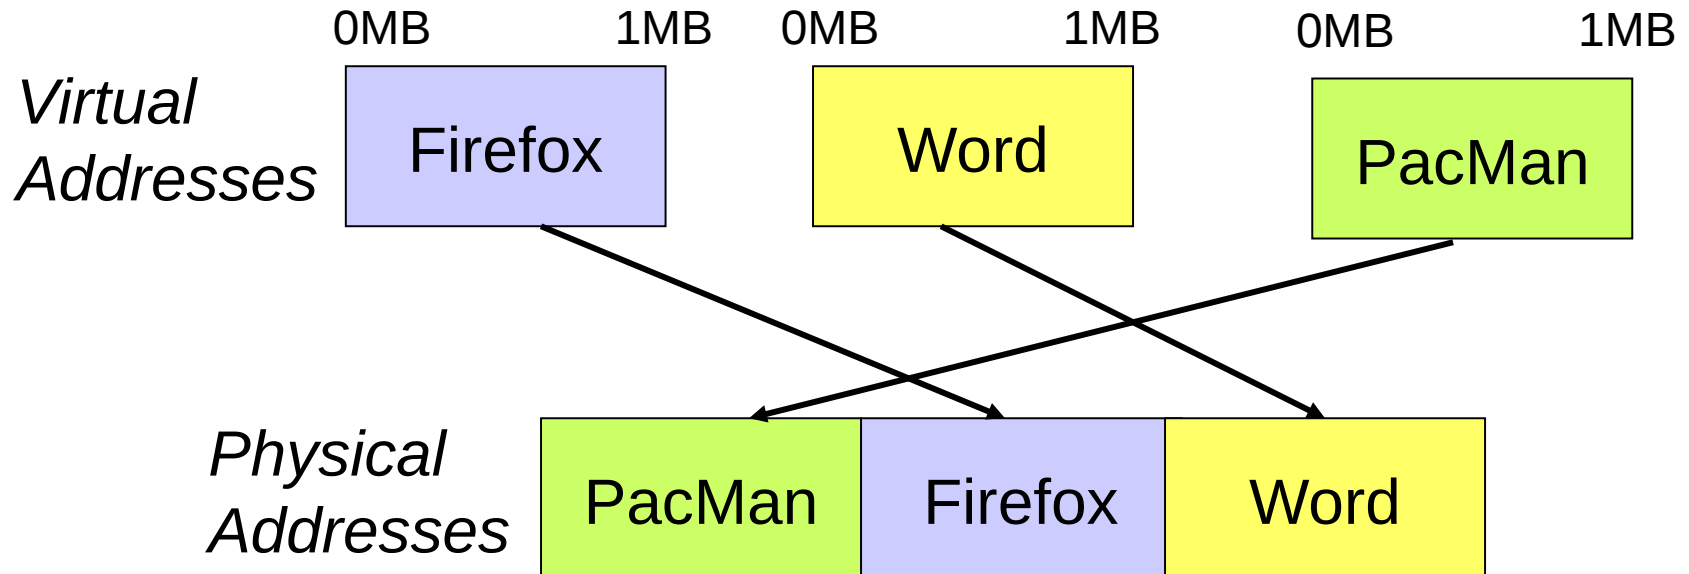


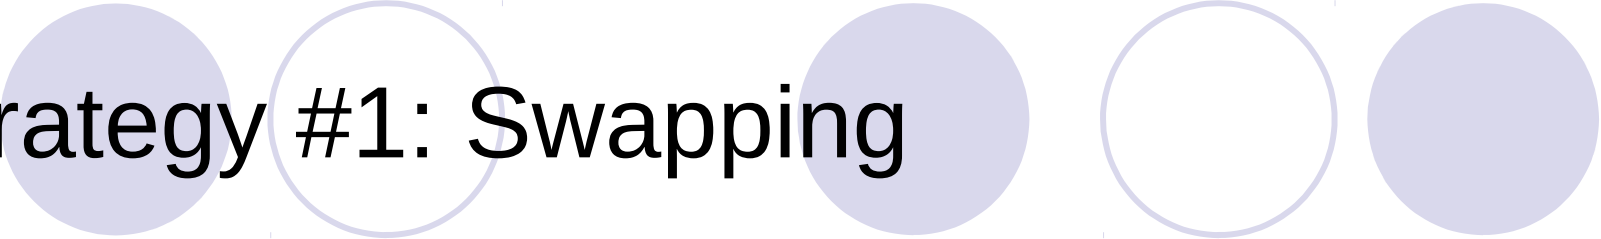
0 MB



# Virtual Addressing

- Add a layer of indirection between user addresses and hardware addresses
- Programmer sees a virtual **address space**, which always starts at zero





# Strategy #1: Swapping

- Only one program occupies memory at once
- On a context switch:
  - Save old program's memory to disk
  - Load next program's memory from disk
- Advantage: very simple
- Disadvantage: disk is exceedingly slow!

# Why Use Swapping?

- No Hardware Support
  - MIT's CTSS operating system (1961) was built this way
    - Memory was so small that only one job would fit!
- Long-lived batch jobs
  - Job switching costs become insignificant as the quantum grows large



# Strategy #2: Fixed partitions

- Physical memory is broken up into fixed partitions
  - All partitions are the same size
  - Partitions never change
- Advantages
  - Simple

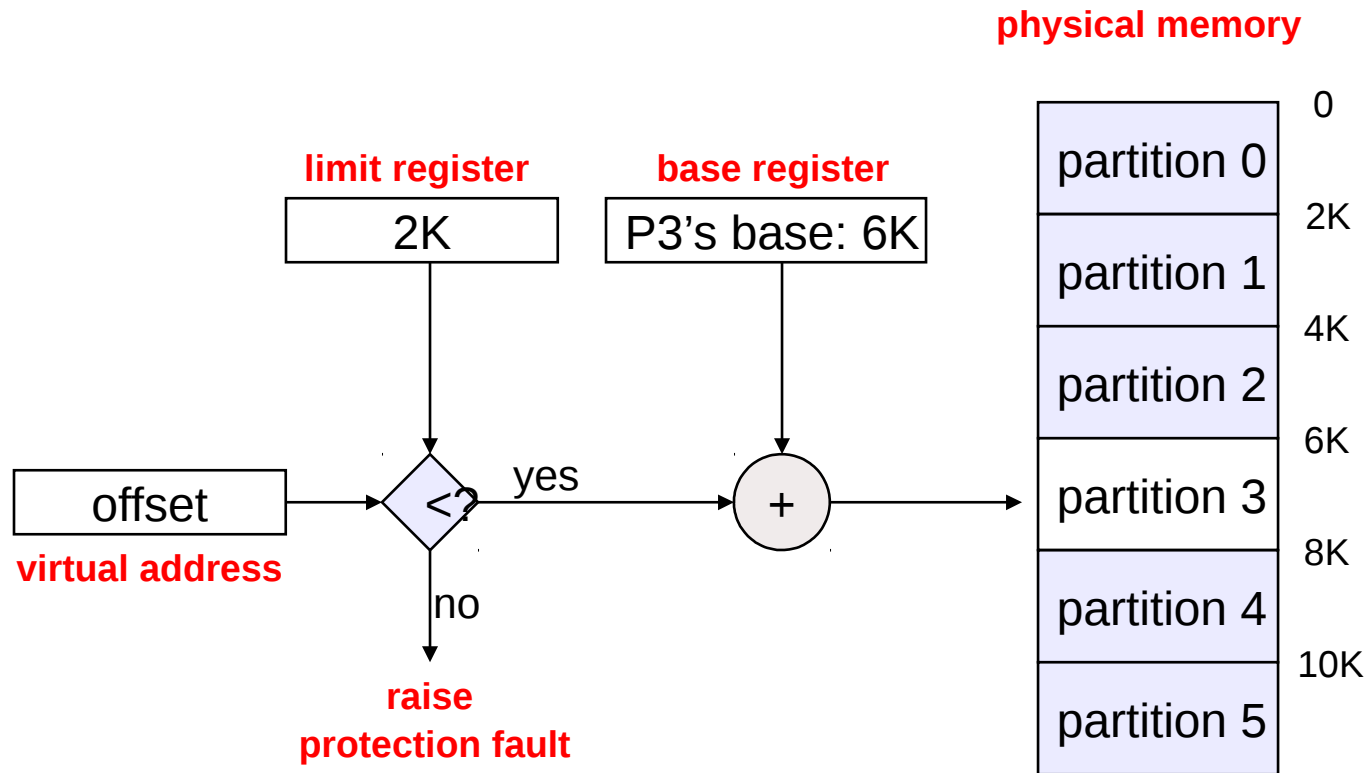


# Disadvantages of Fixed Partitions

- **Internal fragmentation**: memory in a partition not used by its owning process isn't available to other processes
- **Partition size** problem: no one size is appropriate for all processes
  - Tradeoff between fragmentation and accommodating large programs

# Implementing Fixed Partitions

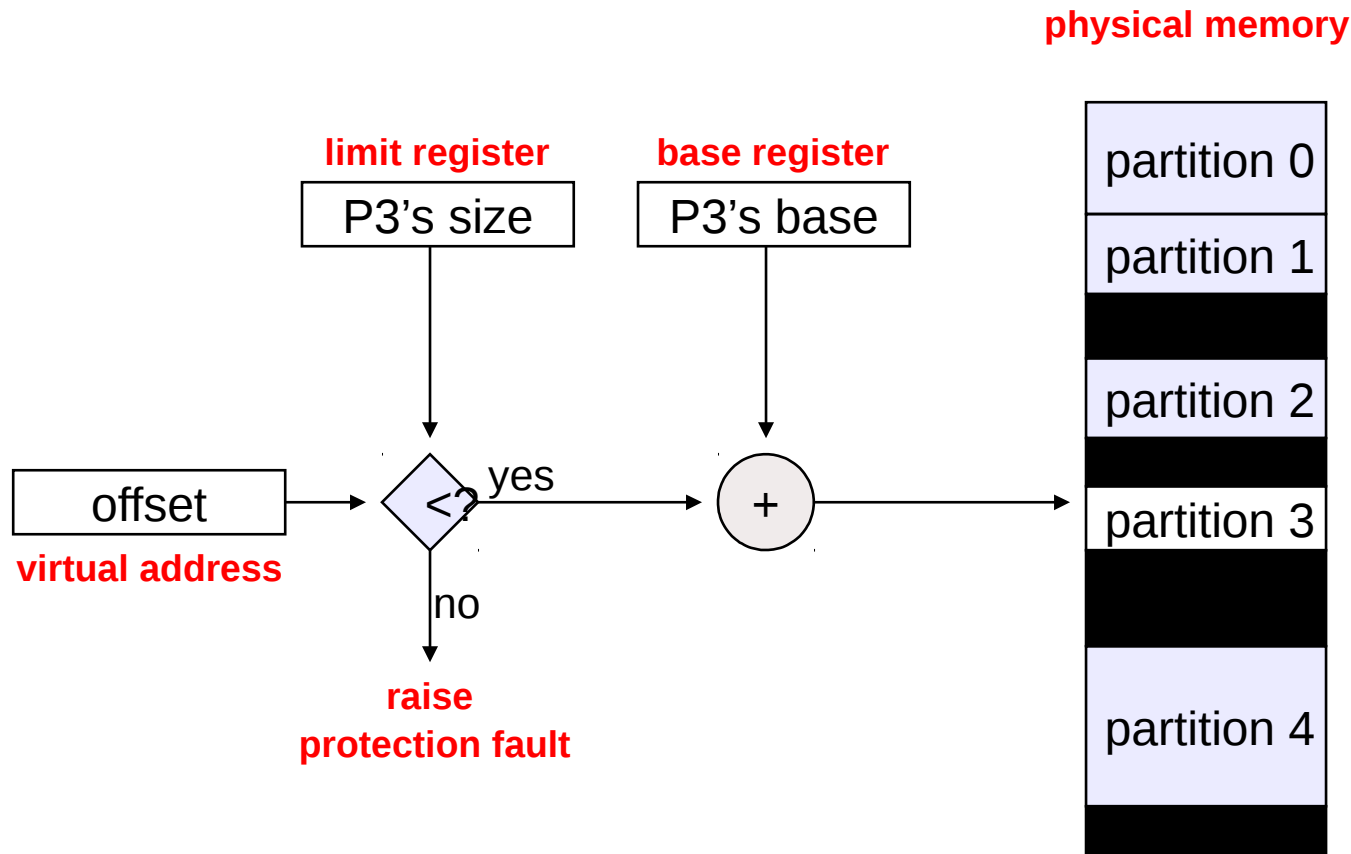
Requires hardware-supported base and limit registers



# Strategy #3: Variable Partitions

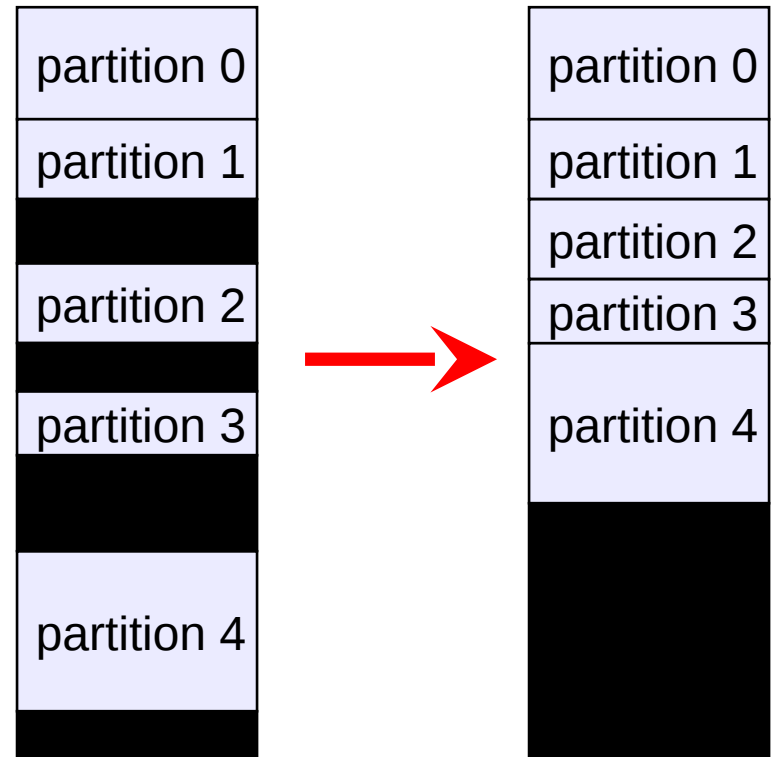
- Obvious next step: physical memory is broken up into variable-sized partitions
- Advantages
  - No internal fragmentation
    - Simply allocate partition size to be just big enough for process
- Problems
  - We must know in advance the program size
  - External fragmentation
    - As we load and unload jobs, holes are left scattered throughout physical memory

# Mechanics of Variable Partitions




# Dealing with External Fragmentation

- Swap a program out
- Re-load it, adjacent to another
- Adjust its base register
- “Lather, rinse, repeat”
- Ugh




# Strategy #4: Paging

- Use fixed-size units of memory (called pages) for both virtual and physical memory



# Paging Advantages



- Paging reduces internal fragmentation
  - How?
- Paging eliminates external fragmentation
  - How?
- Disadvantages?

# Address translation

- A virtual address has two parts: **virtual page number** and **offset**



- Address translation only applies to the virtual page number
  - Why?
- Virtual page number (VPN) is index into a **page table**
  - Page table entry contains **page frame number** (PFN)
  - Physical address is PFN::offset

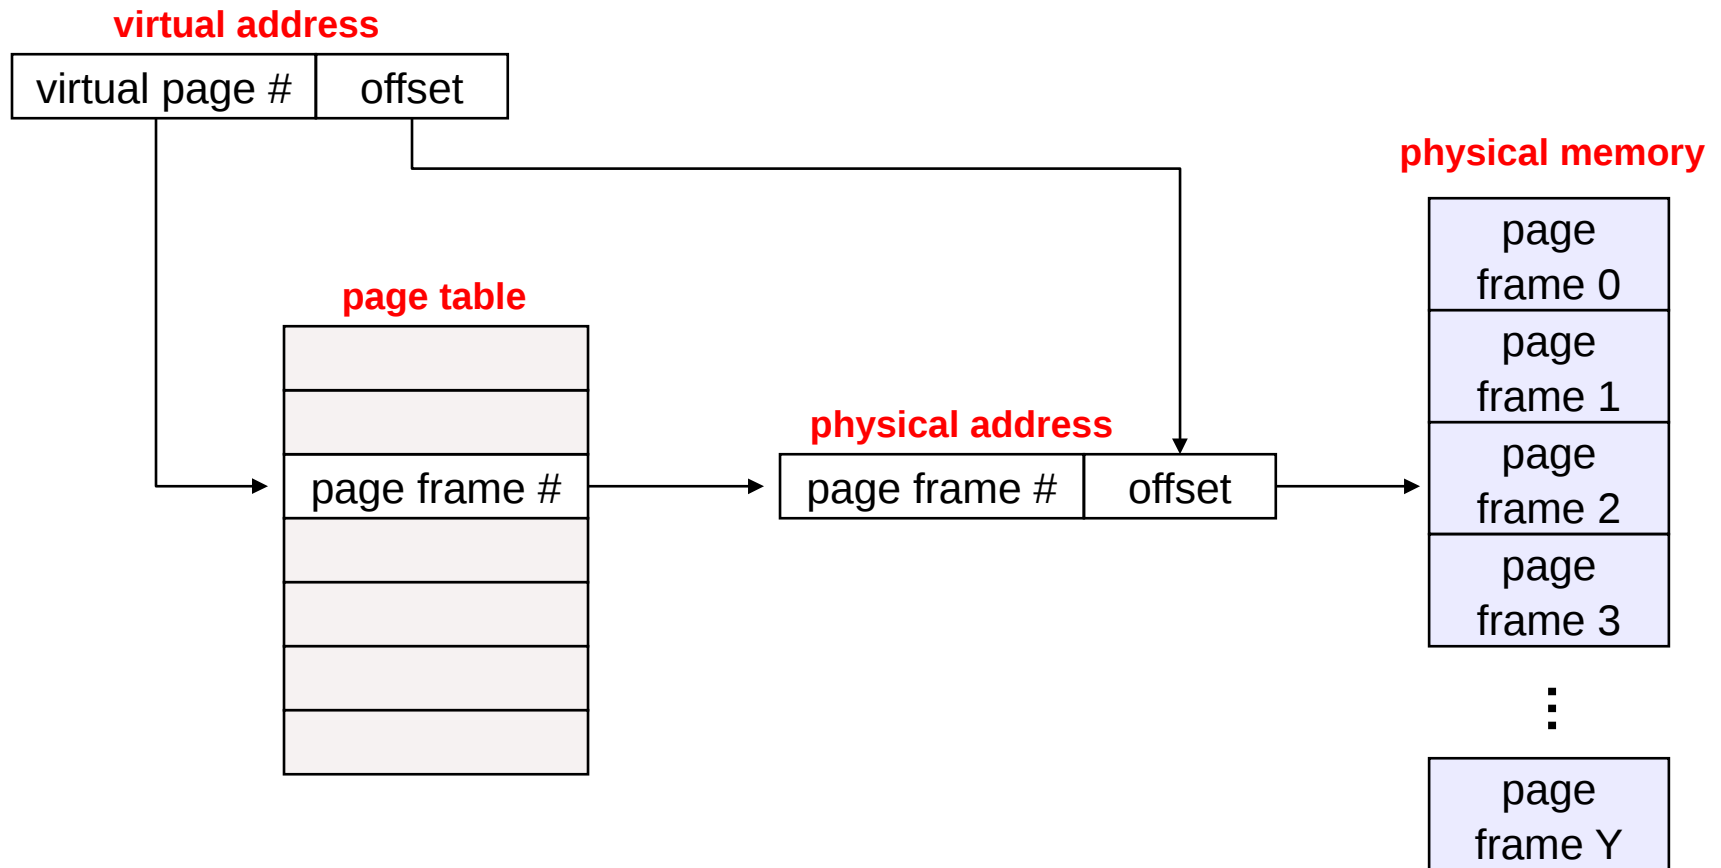


# Page Tables



- Managed by the OS
- Map a virtual page number (VPN) to a page frame number (PFN)
  - VPN is simply an index into the page table
- One **page table entry** (PTE) per page in virtual address space
  - i.e., one PTE per VPN


# Mechanics of address translation



# Example of address translation

- Assume 32 bit addresses
  - assume page size is 4KB (4096 bytes, or  $2^{12}$  bytes)
  - VPN is 20 bits long ( $2^{20}$  VPNs), offset is 12 bits long
- Let's translate virtual address 0x13325328
  - VPN is 0x13325, and offset is 0x328
  - assume page table entry 0x13325 contains value 0x03004
    - page frame number is 0x03004
    - VPN 0x13325 maps to PFN 0x03004
  - physical address = PFN::offset = 0x03004328

# Page Table Entries (PTEs)




- PTE's control address mappings
  - The **page frame number** indicates the physical page
  - The **valid bit** says whether or not the PTE can be used
    - says whether or not a virtual address is valid
  - The **referenced bit** says whether the page has been accessed recently
  - The **modified bit** says whether or not the page is dirty
    - it is set when a write to the page has occurred
  - The **protection bits** control which operations are allowed
    - read, write, execute

# Paging Issues (Stay Tuned...)

- How to make it fast?
  - Accessing the page table on each memory reference is not workable
- How to deal with memory scarcity
  - Virtual memory
- How do we control the memory overhead of page tables?
  - Need one PTE per page in virtual address space
    - 32 bit AS with 4KB pages =  $2^{20}$  PTEs = 1,048,576 PTEs
    - 4 bytes/PTE = **4MB per page table**
  - OS's typically have separate page tables per process
    - 25 processes = 100MB of page tables

# Two-level page tables

- Problem: page tables can be too large
  - $2^{32}$  bytes in 4KB pages need 1 million PTEs
- Solution: use multi-level page tables
  - “Page size” in first page table is large (megabytes)
  - PTE marked invalid in first page table needs no 2nd level page table
- 1st level page table has pointers to 2nd level page tables
- 2nd level page table has actual physical page numbers in it





# More on two-level page tables

- Tradeoffs between 1st and 2nd level page table sizes
  - Total number of bits indexing 1st and 2nd level table is constant for a given page size and logical address length
  - Tradeoff between number of bits indexing 1st and number indexing 2nd level tables
    - More bits in 1st level: fine granularity at 2nd level
    - Fewer bits in 1st level: maybe less wasted space?

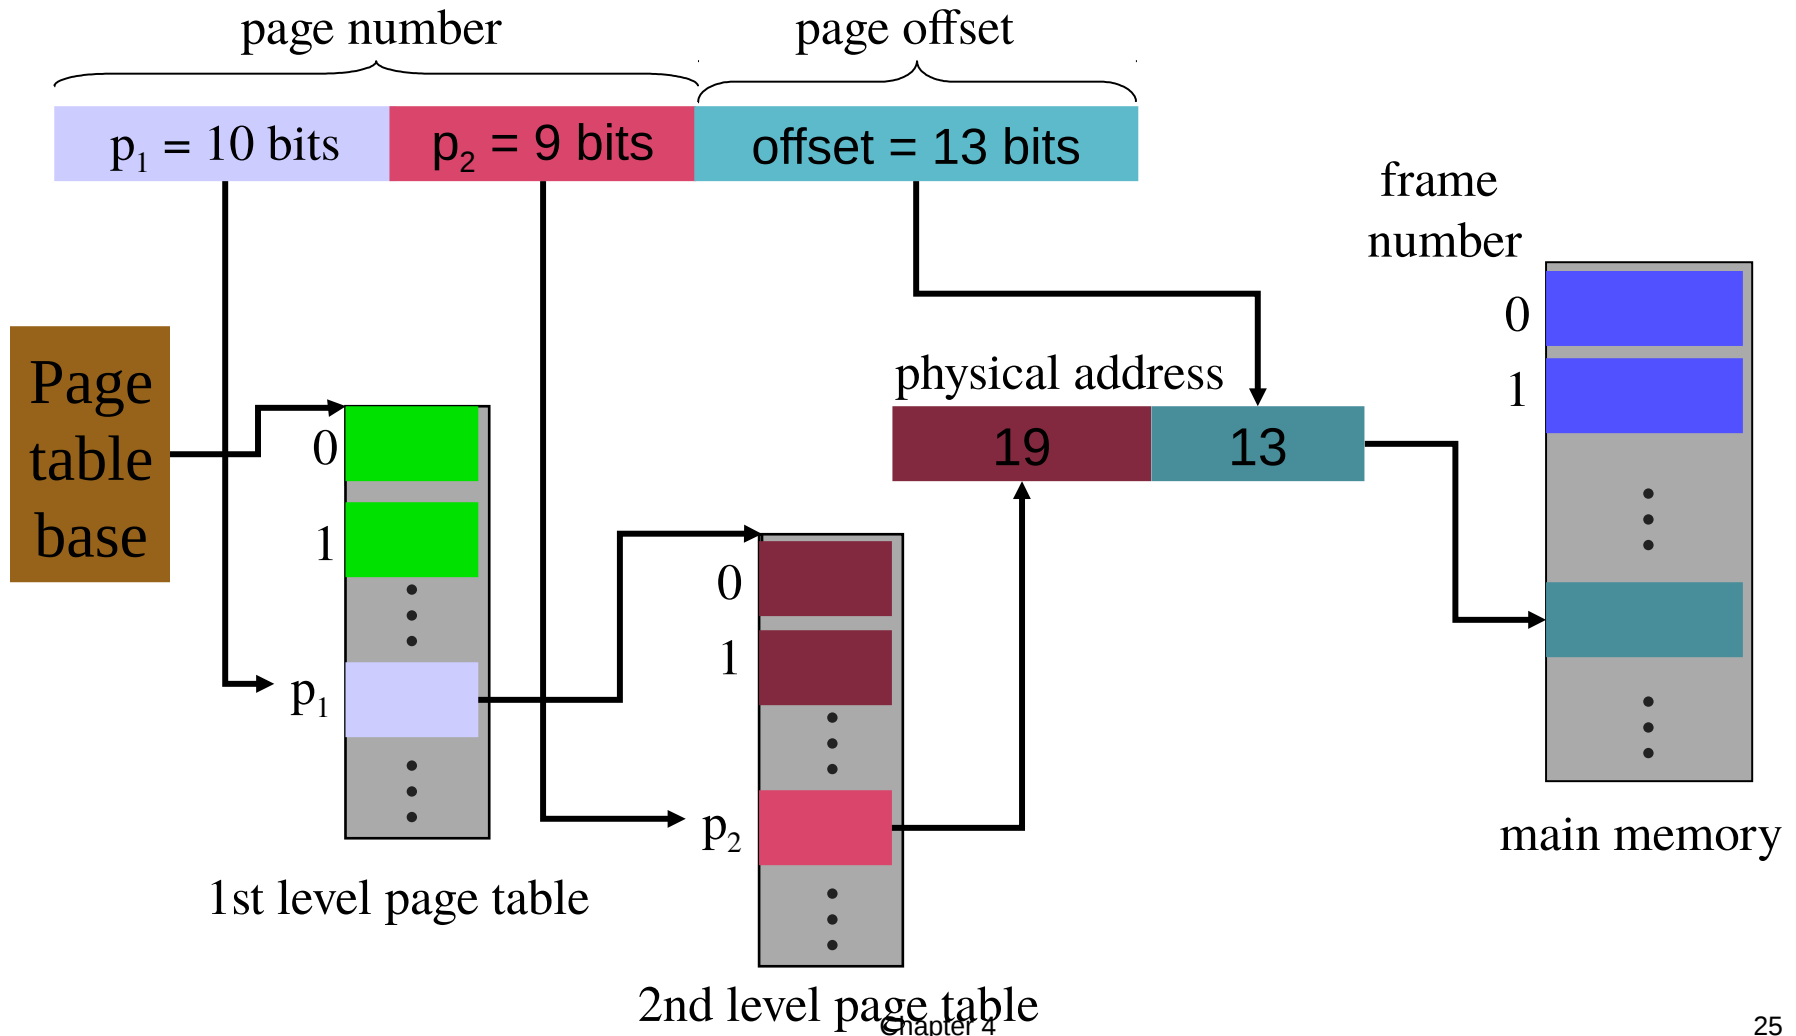
# Two-level paging: example

- System characteristics
  - 8 KB pages
  - 32-bit logical address divided into 13 bit page offset, 19 bit page number
- Page number divided into:
  - 10 bit page number
  - 9 bit page offset
- Logical address looks like this:
  - $p_1$  is an index into the 1st level page table
  - $p_2$  is an index into the 2nd level page table pointed to by  $p_1$






# 2-level address translation example




# Strategy #5: Segmentation


- Instead of a flat address space, programmers see a collection of segments



- Virtual address = segment #, offset



# Why Segments?



- Facilitates sharing and re-use
  - Unlike a page, a segment is a **logical** entity
    - Segments can have arbitrary size