

RR schedulers with Multi-level queuing and Multi-processor scheduling

10-Aug-2016

CS303

Autumn 2016

Motivation





- Process types:
 - CPU-bound Vs I/O-bound
- There are several sched. Algos.: (we have seen...)
 - FCFS, SJF, RR, SJRT
 - SJF/SJRT provide better avg. wait time and Throughput

Can RR be made to exploit the nature of a process?

- RR can be adapted to expedite IO-bound process than a CPU-bound process
 - Consequently RR also exhibits improved
 - avg. wait time and throughput etc.
- How?
 - With Multi-level queuing we can achieve this improvement

Multilevel queuing (1/2)

- Ready queue in RR can be conceptualised as a multi-level queue
 - consisting of multiple queues with different time-quanta
 - The queue with lower time-quanta is given more priority
 - Within the Queue Q_i , the processes go RR with timequanta T_i
- Multi-level queue


		TIME QUANTA	CBT interval $CI[i]$
Q1		$T1$	$(0, T1]$
Q2		$T2$	$(T1, T2]$
Q3		$T3$	$(T2, T3]$
Q4		$T4$	$(T3, T4]$

Such that $T1 < T2 < T3 < T4$

And Priority: $\pi(Q1) > \pi(Q2) > \pi(Q3) > \pi(Q4)$

Multilevel queuing (2/2)

- Within every level/queue all the processes are scheduled in RR with corresponding time-quanta




Such that $T1 < T2 < T3 < T4$

And Priority is $\pi(Q1) > \pi(Q2) > \pi(Q3) > \pi(Q4)$

Multilevel queuing: Illustration (1/5)


- Consider process P1 with.CBT 25 milli-Seconds(mSec)
 - There is a 4-level queue with timequantas:
 - 2mS, 5mS 10mS and 15mS
 - Let arrival time be 0mS
- Completed Burst Time:0mS



Multilevel queuing: Illustration (2/5)


- As per the algo., P1 is placed in Q1 with TQ of 2mS
- Waits for its turn on RR Scheduler of Q1

Completed Burst Time:0mS



Multilevel queuing: Illustration (2/5)


- As per the algo., P1 is placed in Q1 with TQ of 2mS
- Waits for its turn on RR Scheduler of Q1
- RR scheduler of Q1 dispatches P1, P1 starts running
Completed Burst Time:0mS



Multilevel queuing: Illustration (3/5)

- P1 executes for 2mS; execution does not finish
- Gets preempted

Completed Burst Time: 2mS




Multilevel queuing: Illustration (4/5)

- Since execution not finished
- It is moved down to level-2

Completed Burst Time: 2mS


P1 is pushed
Down from Q1
to Q2



Multilevel queuing: Illustration (5)

- With 2mS the execution is not finished
- It is moved down to level-2

Completed Burst Time:2mS



Multilevel queuing: Illustration (6/5)


- With 2mS the execution is not finished
- It is moved down to level-2
- **In Q2, on getting its turn in RR sched, starts execution**

Completed Burst Time:2mS

P1 gets dispatched
and starts to run
from Q2


TIME QUANTA

Q1




2mS

Q2



5mS

Q3



10mS

Q4




15mS

Multilevel queuing: Illustration (7/5)

- With 2mS the execution is not finished
- It is moved down to level-2
- In Q2, on getting its turn in RR sched, starts execution
- **P1 exhausts its quanta of 5mS in Q2; gets preempted**

Completed Burst Time: 7mS

P1's Quanta
exhausted
in Q2




Multilevel queuing: Illustration (8/5)

In Q2, on getting its turn in RR sched, starts execution

- Process exhausts its quanta of 5mS in Q2
- P1 moves out to Q3

P pushed
Down from Q2
to Q3

Completed Burst Time: 7mS




Multilevel queuing: Illustration (9/5)

Process exhausts its quanta of 5mS in Q2

- P1 moves out to Q3
- P1 in Q3
- P1 waits for its turn on RR sched of Q3

Completed Burst Time: 7mS



Multilevel queuing: Illustration (10/5)


P1 moves out to Q3

- P1 in Q3
- P1 waits for its turn on RR sched of Q3
- P1 is dispatched and start execution

Completed Burst Time: 7mS

P1 begins exec
Post dispatch
From Q3 by its RR sched

TIME QUANTA



Multilevel queuing: Illustration (11)

P1 in Q3


- P1 waits for its turn on RR sched of Q3
- P1 is dispatched and start execution
- P1 exhausts its quanta of 10 mS in Q3, gets preemted

Completed Burst Time:17mS

P1 executes for 10mS
After entering Q3
Still to exit


TIME QUANTA

Q1




2mS

Q2



5mS

Q3



10mS


Q4



15mS

Multilevel queuing: Illustration (11)

- P1 exhausts its quanta of 10 mS in Q3
- P1 is pushed down to Q4




Multilevel queuing: Illustration (11)

P1 exhausts its quanta of 10 mS in Q3

- P1 is pushed down to Q4
- P1 waits from Q4 for its turn in the RR sched. Of Q4

P1 waits in Q4
For its turn in RR
Of Q4

Completed Burst Time:17mS




Multilevel queuing: Illustration (11)

P1 is pushed down to Q4

- P1 waits from Q4 for its turn in the RR sched. Of Q4
- RR Scheduler of Q4 when turn for P1 dispatches it, P1 execution begins

P1 gets dispatched
by RR scheduler of
Of Q4

Completed Burst Time:17mS




Multilevel queuing: Illustration (11)

RR Scheduler of Q4 when turn for P1 dispatches it, P1 execution begins

- Runs for another 8mS and exits voluntarily
- The queue-level-association is done ONLY after EXIT

Completed Burst Time:17mS

P1 runs for another 8 mS and exits



Proc-Queue Association

process	Queue


Multilevel queuing: Illustration (11)

P1 waits from Q4 for its turn in the RR sched. Of Q4

- RR Scheduler of Q4 when turn for P1 dispatches it, P1 execution begins
- Runs for another 8mS and exits voluntarily

Completed Burst Time:17mS

Process P1 – Q4
Association is done



Mapping Table

process	Queue
P1	Q4


On exit, only Queue-mapping
gets fixed

Multilevel queuing: Illustration (11)

- Runs for another 8mS and exits voluntarily
- NEXT time onwards the P1 **always** enters Q4

P1 runs for another
8 mS and exits

Completed Burst Time:17mS



Mapping Table

process	Queue
P1	Q4

Process allocation to each queue-level

- **Whenever a process becomes ready the process is allocated to a queue as following:**

1. place it in the top-level queue Q_i (initialise $i=1$)

2. **IF** it's (remaining) CPU-Burst within T_i

THEN It is associated to this queue (permanently)

 // all processes within a queue-level get RR-scheduled

 //in ALL subsequent rounds it enters Q_1 always

ELSE IF ($i < n$) {

$i++$; //move it to next-level

 move process to Q_i ;

GOTO Step2

 }

ELSE associate it to Q_n ;

3. From next time a process becoming RDY, enters its associated Q_i

4. If queues above Q_i are empty, processes in Q_i are scheduled in RR with quanta T_i

This association is permanent: a process always enters that Q

Process allocation to each queue-level

- Whenever a process becomes ready the process is allocated to a queue as following:

1. place it in the top-level queue Q_i (initialise $i=1$)

2. IF it's (remaining) CPU-Burst within T_i

THEN It is associated to this queue (permanently)

 // all processes within a queue-level get RR-scheduled

 //in ALL subsequent rounds it enters Q_1 always

ELSE IF ($i < n$) {

$i++$; //move it to next-level

 move process to Q_i ;

GOTO Step2

 }

ELSE associate it to Q_n ;

3. From next time a process becoming RDY, enters its associated Q_i

4. If queues above Q_i are empty, processes in Q_i are scheduled in RR with quanta T_i

This association is permanent: a process always enters that Q

Process allocation to each queue-level

- Whenever a process becomes ready the process is allocated to a queue as following:

1. place it in the top-level queue Q_i (initialise $i=1$)

- 2. IF** it's (remaining) CPU-Burst within T_i

THEN It is associated to this queue (permanently)

 // all processes within a queue-level get RR-scheduled

 //in ALL subsequent rounds it enters Q_1 always

ELSE IF ($i < n$) {

$i++$; //move it to next-level

 move process to Q_i ;

GOTO Step2

 }

ELSE associate it to Q_n ;

3. From next time a process becoming RDY, enters its associated Q_i

4. If queues above Q_i are empty, processes in Q_i are scheduled in RR with quanta T_i

This association is permanent: a process always enters that Q

Process allocation to each queue-level

- Whenever a process becomes ready the process is allocated to a queue as following:

1. place it in the top-level queue Q_i (initialise $i=1$)

- 2. IF** it's (remaining) CPU-Burst within T_i

THEN It is associated to this queue (permanently)

// all processes within a queue-level get RR-scheduled

//in ALL subsequent rounds it enters Q_1 always

ELSE IF ($i < n$) {

$i++$; //move it to next-level

move process to Q_i ;

GOTO Step2

}

ELSE associate it to Q_n ;

3. From next time a process becoming RDY, enters its associated Q_i

4. If queues above Q_i are empty, processes in Q_i are scheduled in RR with quanta T_i

This association is permanent: a process always enters that Q

Process allocation to each queue-level

- Whenever a process becomes ready the process is allocated to a queue as following:

1. place it in the top-level queue Q_i (initialise $i=1$)

- 2. IF** it's (remaining) CPU-Burst within T_i

THEN It is associated to this queue (permanently)

 // all processes within a queue-level get RR-scheduled

 //in ALL subsequent rounds it enters Q_1 always

ELSE IF ($i < n$) {

$i++$; //move it to next-level

 move process to Q_i ;

GOTO Step2

 }

ELSE associate it to Q_n ;

3. From next time a process becoming RDY, enters its associated Q_i

4. If queues above Q_i are empty, processes in Q_i are scheduled in RR with quanta T_i

This association is permanent: a process always enters that Q

Process allocation to each queue-level

- Whenever a process becomes ready the process is allocated to a queue as following:

1. place it in the top-level queue Q_i (initialise $i=1$)

- 2. IF** it's (remaining) CPU-Burst within T_i

THEN It is associated to this queue (permanently)

 // all processes within a queue-level get RR-scheduled

 //in ALL subsequent rounds it enters Q_1 always

ELSE IF ($i < n$) {

$i++$; //move it to next-level

 move process to Q_i ;

GOTO Step2

 }

ELSE associate it to Q_n ;

3. From next time a process becoming RDY, enters its associated Q_i

4. If queues above Q_i are empty, processes in Q_i are scheduled in RR with quanta T_i

This association is permanent: a process always enters that Q

Draw backs of static multi-level-Q RR sched.

- ML-RR scheduler is
 - insensitive to processes that change their nature from round to round
 - i.e. IO-bound becoming CPU-bound or vice versa
 - Prediction cannot be so quick
- Solution: Make association NOT fixed for more than one round
 - i.e. dynamic mapping
 - Consider the association only for next round, then evaluate again and decide the future association
 - Consequently
 - IF a process shifts from IO-bound to CPU-bound we migrate it downwards
 - ELSE IF a process shifts from CPU-bound to IO -bound we migrate it upwards sadf
- In dynamic version, we can have:
 - Feed-back multi-level queue RR scheduling

The whole motive is to improve efficiency from wait-time, TAT, CPU-utilisation, etc.

Feed-back Multi-level RR Scheduling

- Post first-time association:
 1. CPU-burst (CBT_p) of each process P_p finishing in Q_i is recorded // I is the queue-level
 2. IF CBT_p in CI[j], for some level 'j' above level i:
THEN push the process to that level 'j' upwards //FEED-BACK
ELSE IF ExecutionTime of p 'ET_p' exceeds T_i
THEN move p to queue one-level downwards //FEED-FORWARD /
// ELSE otherwise the process sticks to the same queue

THEREFORE, any process on changing its nature is sensed by this RR ML-scheduling
THUS, exploits this change by juggling the processes up or down!


Accordingly we have connections between the heads and tails of various levels.

But, a caution: more context switches means reduction in efficiency, (RECALL the tutorial problem on GANTT CHARTS), more context switches means reduced CPU-utilisation, and increased wait time and TAT !!!!!

Such checking of times and comparison is costly, we can't afford to loose the efficiency by investing more processor time on scheduling!!!!


Feed-back Multi-level RR Scheduling

- > any process on changing its nature is sensed by this RR ML-scheduling
- > THUS juggling the processes up or down accordingly improves efficiency
- > Accordingly we have connections between the heads and tails of various levels.



Feed-back Multi-level RR Scheduling

- > Feed-forward: IO-bound becoming CPU-bound
- > Several configs are possible connecting tails of lower levels to heads of multiple levels up



Feed-back Multi-level RR Scheduling

CAUTION

- But, a caution:
 - more context switches means reduction in efficiency, (RECALL the tutorial problem on GANTT CHARTS), more context switches means reduced CPU-utilisation, and increased wait time and TAT !!!!!
- Such checking of times and comparison is costly, we can't afford to loose the efficiency by investing more processor time on scheduling!!!!