


Memory Management (MFT)

3 November 2017


CS303

Autumn 2017

Memory Management

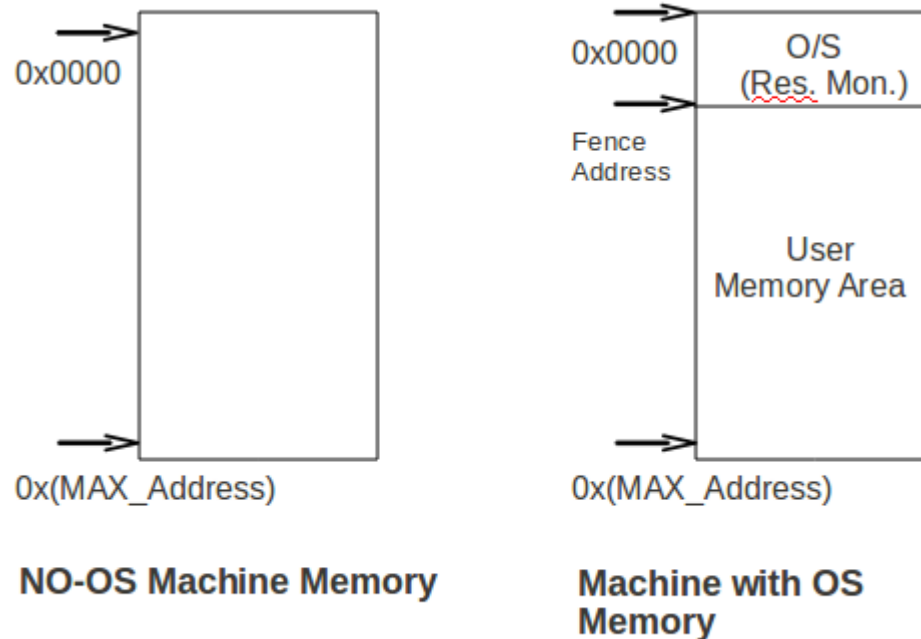


NO-OS Machine Memory



Machine with OS Memory

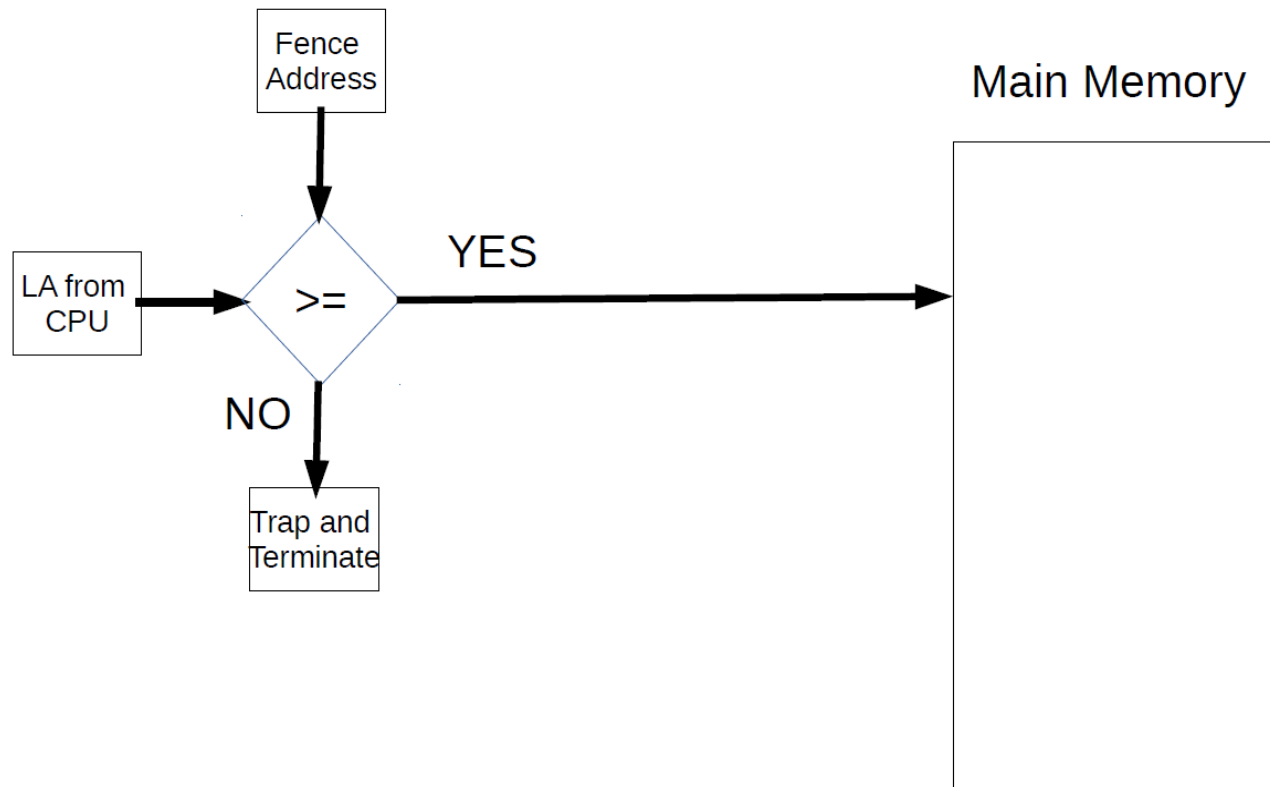
Memory Management



Fence Address (FA): stored in a register
Either Fixed in HW or Editable

1. On every address generated by CPU from a user process, it needs to be compared with the FA
 - this check can be implemented in
 - HW: hardware comparator 'Efficient'
 - SW: software instruction for comparing


Memory Management



Fence Address (FA): stored in a register
Either Fixed in HW or Editable



1. On every address generated by CPU from a user process, it needs to be compared with the Fence Address
 - this check can be implemented in
 - HW: hardware comparator 'Efficient'
 - SW: software instruction for comparing

Memory Management




1. Fence Address (FA): stored in a register
Either Fixed in HW or Editable
2. The checking requires a comparator, which could be implemented in:
 - HW: hardware comparator 'Efficient'
 - SW: software instruction for comparing

Simple Memory Model: Res. Montior + User Area



1. Single Process gets loaded into the MM from the SS/HD

Simple Memory Model: Res. Monitor + Use Area




1. Single Process gets loaded into the MM from the SS/HD
2. When other process needs to run, first the current running process has to be moved-out onto the secondary storage

DRAWBACKS:

- The CPU would be idle during this IO operation, so inefficient
- Also, the degree of Multiprogramming is just 1.


Multi-user System with OS



OBSERVE: It involves a kind of IO operation – as Secondary Storage is also a kind of IO device

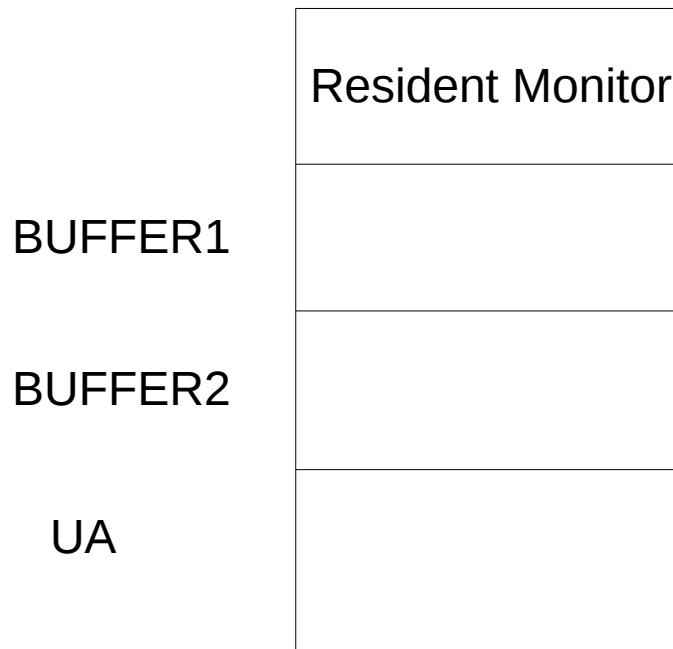
1. Sec. Storage is slower than MM access rate
2. The changes in var values etc. need to be changed in the SM as well

Multi-user System with OS



1. On demand, process is moved-out onto Sec. Storage
2. It involves IO operation
 - Sec. Storage is slower than MM access rate
3. The CPU would wait in this IO operation, so inefficient
4. Drawback back: Degree of Multiprogramming is just 1.


Partitioned Memory Model: Multiple Buffer Partitions + Active User Area



Partitions: Set of non-overlapping memory regions called partitions

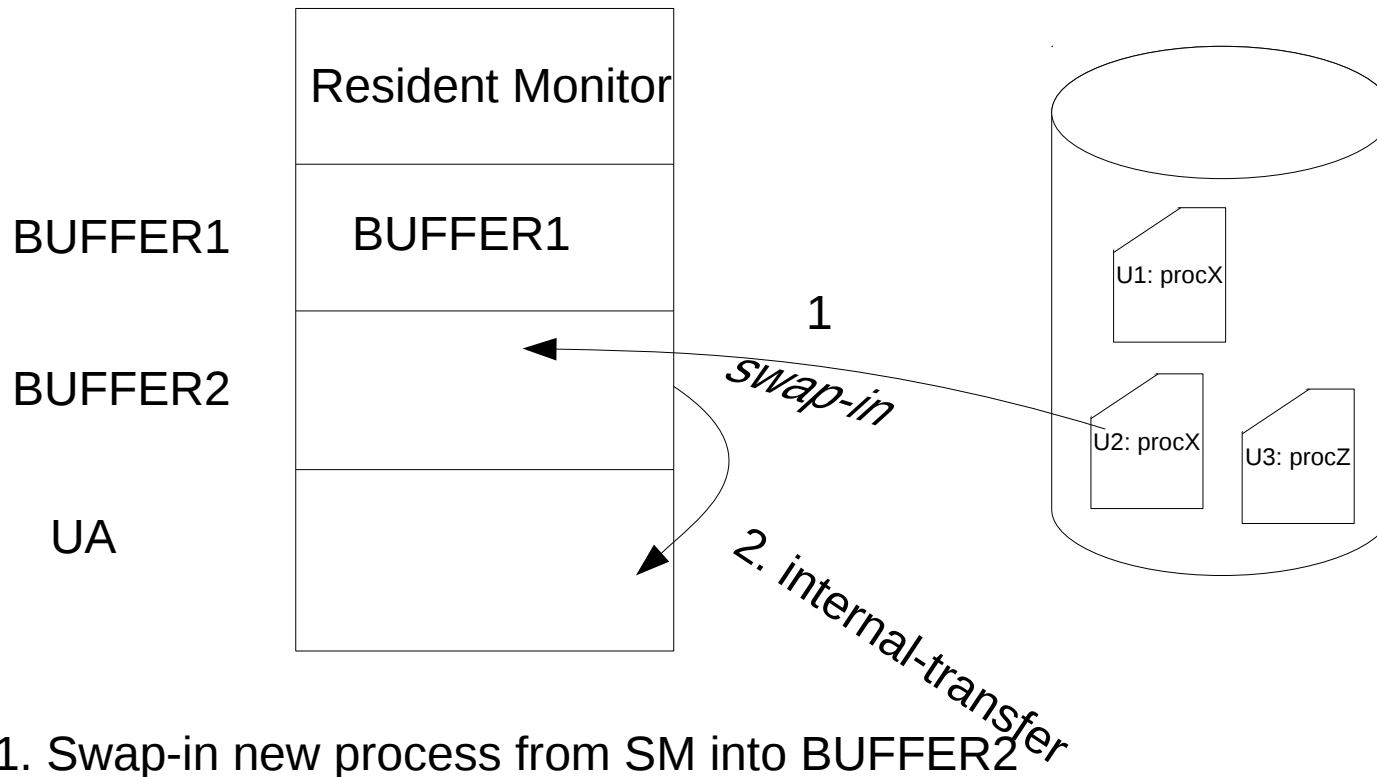
- I .Partition the UA of MM into
BUFFER1, BUFFER2, ..., User Area
- II. This makes possible for:
 1. Swap-in new process from SM into BUFFER2
 2. **Move the new process from BUFFER2 to UA**
 3. **Move Active process from UA to BUFFER1**
 4. Swap-out process from BUFFER1 onto SM
- III. Multiple progs can be loaded into buffers simultaneously

Partitioned Memory Model: Multiple Buffer Partitions + Active User Area



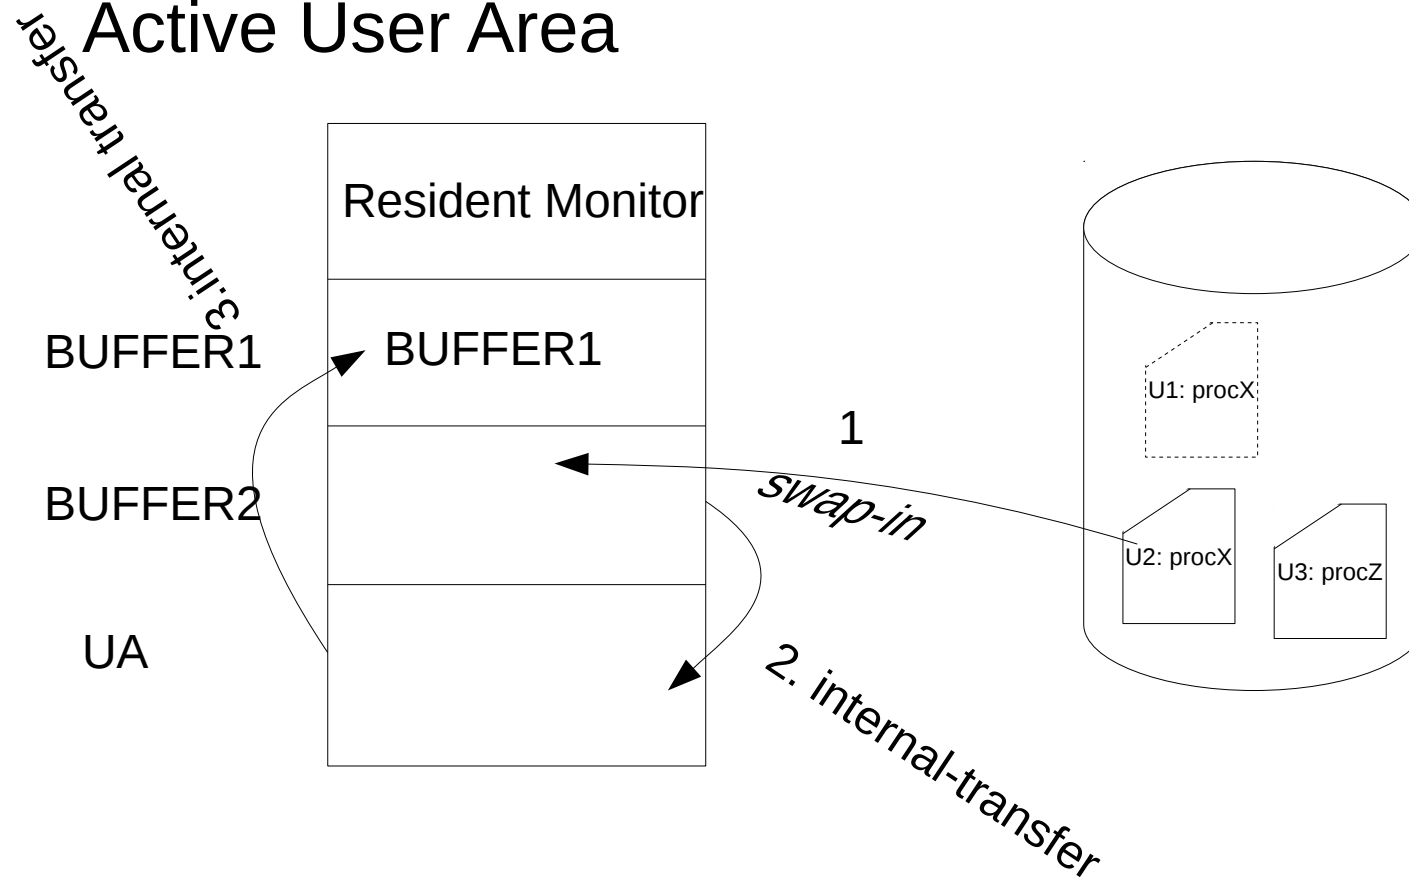
1. Swap-in new process from SM into BUFFER2

Partitioned Memory Model: Multiple Buffer Partitions + Active User Area



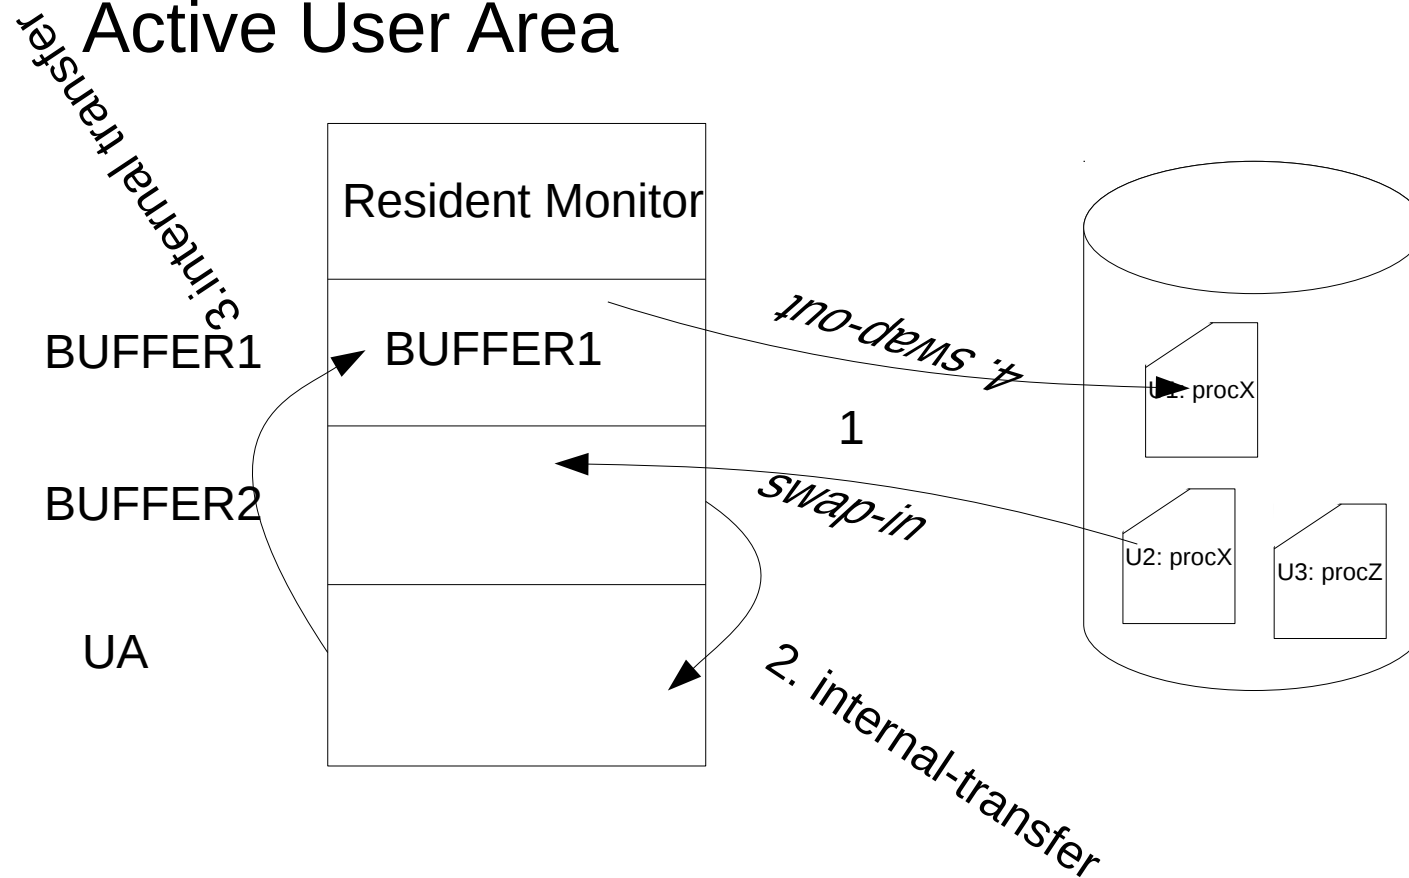
1. Swap-in new process from SM into BUFFER2
2. **Move the new process from BUFFER2 to UA**

Partitioned Memory Model: Multiple Buffer Partitions + Active User Area



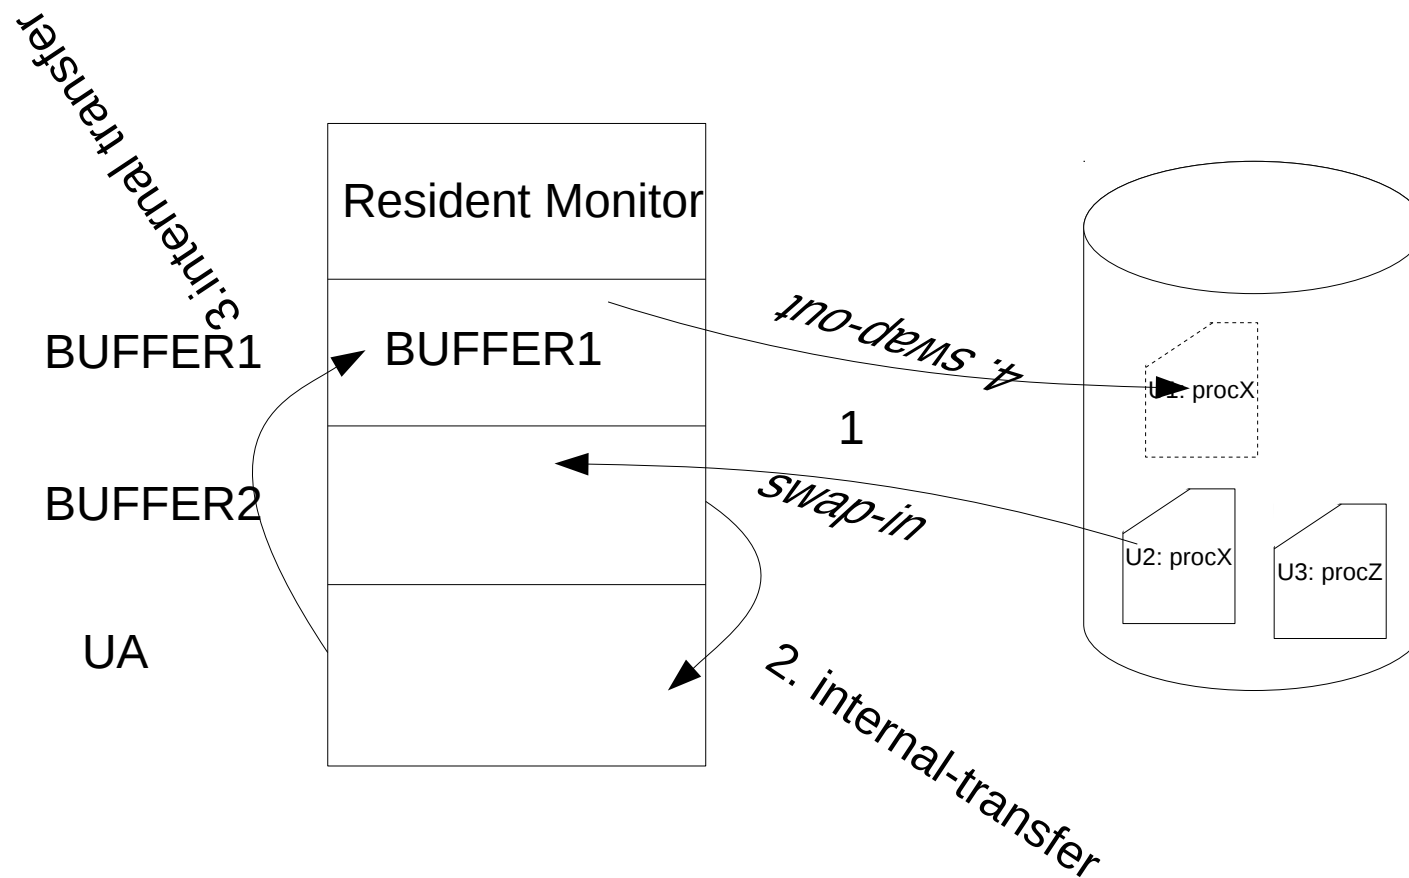
1. Swap-in new process from SM into BUFFER2
2. **Move the new process from BUFFER2 to UA**
3. **Move Active process from UA to BUFFER1**

Partitioned Memory Model: Multiple Buffer Partitions + Active User Area



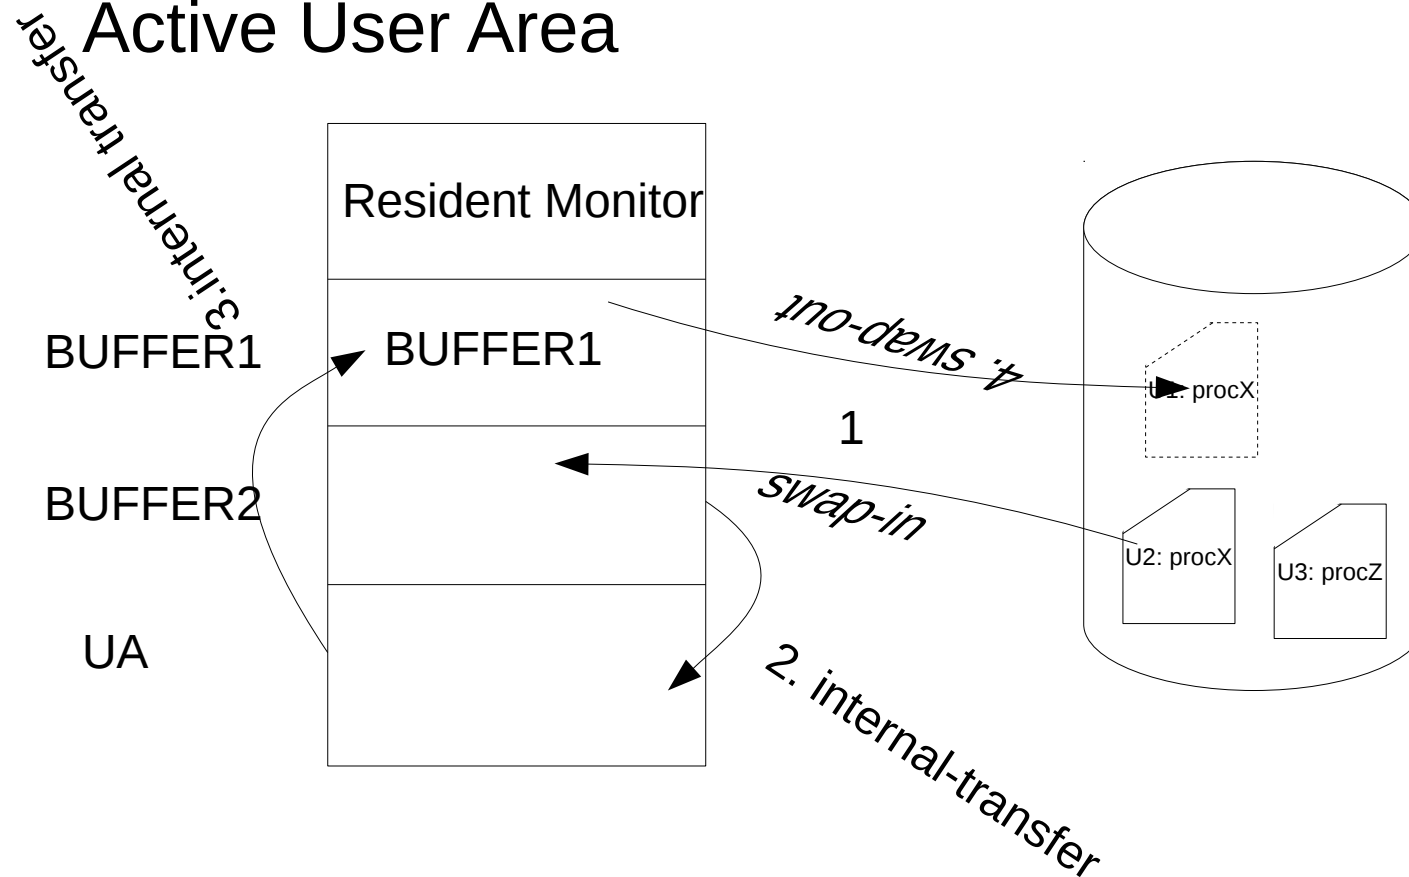
1. Swap-in new process from SM into BUFFER2
2. **Move the new process from BUFFER2 to UA**
3. **Move Active process from UA to BUFFER1**
4. Swap-out process from BUFFER1 onto SM

Multi-partitioned Memory Model




3. Multiple progs can be loaded into buffers simultaneously
- degree of multiprogramming increases

Partitioned Memory Model: Multiple Buffer Partitions + Active User Area



Multiple progs can be loaded into buffers simultaneously
- degree of multiprogramming increases


Multi-Partitioned Memory Model: Fixed number of fixed sized partitions



1 .Every partition is characterised by:
Base Address
Size/Limit

2. Process when loaded gets mapped to one partition

Multi-Partitioned Memory Model: Fixed number of partitions



1. In such multi-partitioned model, memory protection is more complex
- need to make sure one process in a partition can not access address belonging to another partition

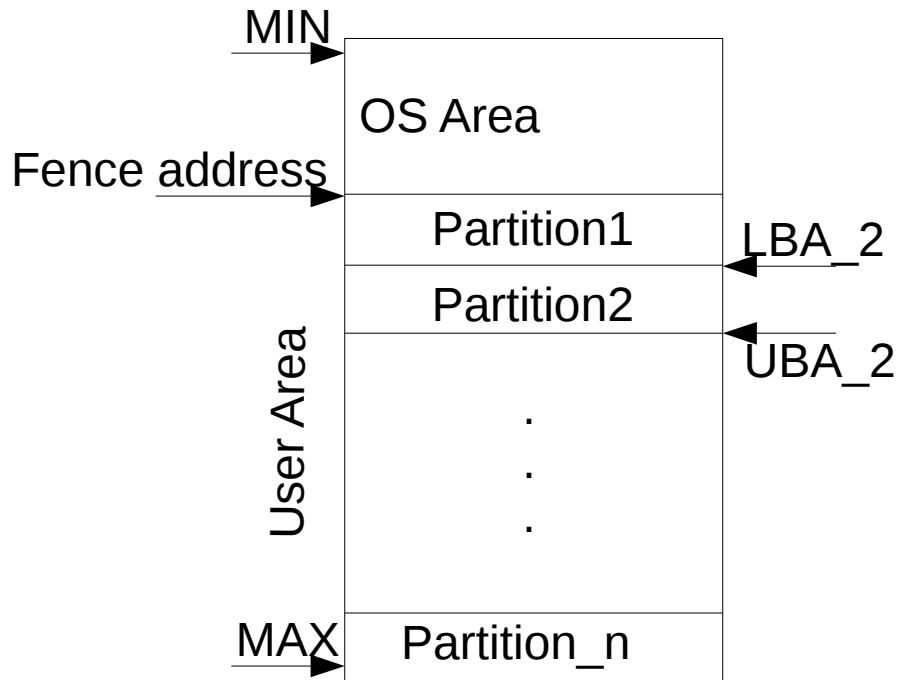
2. Each partition is identified by a contiguous memory addresses between two address bounds;
-Lower Bound Address/Register (LBA)
-Upper Bound Address/Register (UBA)

If size of partition is S_p , then
 $UBR_p = LBR_p + S_p$

3. So every address from a process (belonging to a partition) shall respect these bounds for that partition.

4. LBA is stored in Base Register; while UBA is stored in Limit Register

Multi-Partitioned Memory Model: Fixed number of partitions



1. In such multi-partitioned model, memory protection is more complex
- need to make sure one process in a partition can not access address belonging to another partition

2. Each partition is identified by a contiguous memory addresses between two address bounds;
-Lower Bound Address (LBA)
-Upper Bound Address (UBA)

If size of partition is S_p , then
 $UBA_p = LBA_p + S_p$

3. So every address from a process (belonging to a partition) shall respect these bounds for that partition.

4. Every address shall be greater than 0

Process Partition Mapping Table

| BASE | LIMIT |
|------|-------|
| | |
| | |
| | |
| | |
| | |

Main Memory

| O/S |
|-----|
| |

LA from
CPU

Process Partition Mapping Table

| BASE | LIMIT |
|------|-------|
| | |
| | |
| | |
| | |
| | |



processID

LA from
CPU

Main Memory

O/S

Process Partition Mapping Table


| BASE | LIMIT |
|------|-------|
| | |
| | |
| | |
| | |
| | |



processID

LA from
CPU

Main Memory



Process Partition Mapping Table

| BASE | LIMIT |
|------|-------|
| | |
| | |
| | |
| | |
| | |



processID

limit[i]

base[i]

Main Memory

O/S

LA from
CPU



Process Partition Mapping Table

| BASE | LIMIT |
|------|-------|
| | |
| | |
| | |
| | |

processID

limit[i]

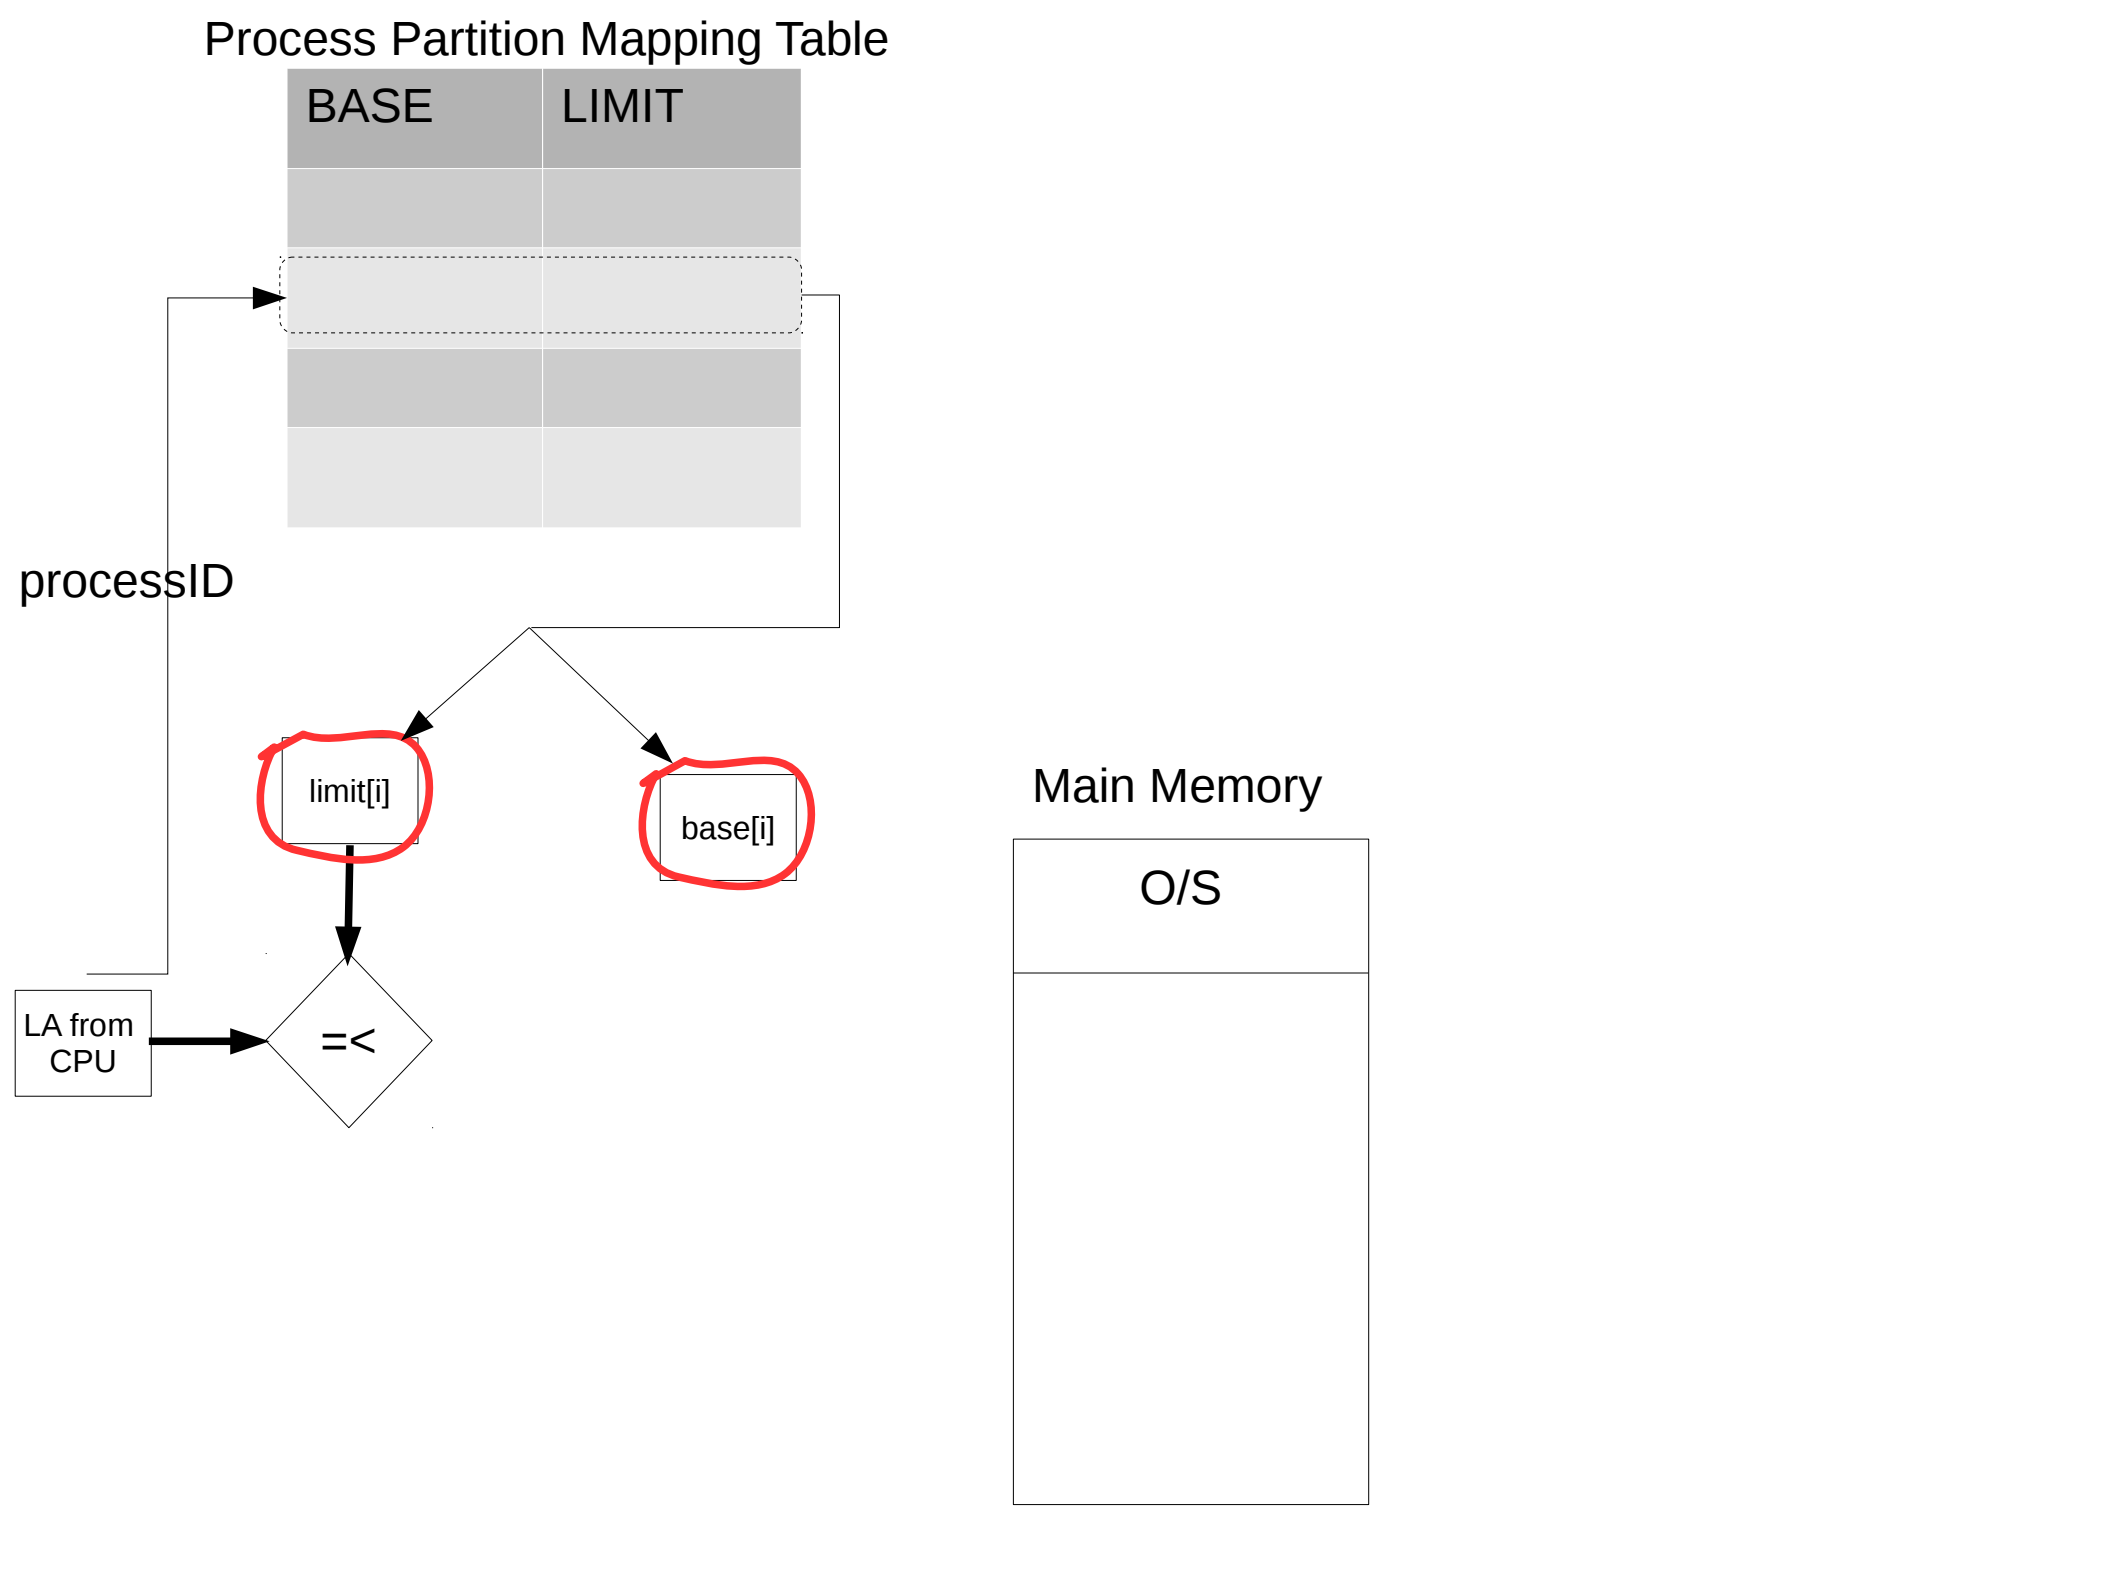
base[i]

Main Memory

O/S

LA from
CPU

=<



Process Partition Mapping Table

| BASE | LIMIT |
|------|-------|
| | |
| | |
| | |
| | |

processID

limit[i]

base[i]

Main Memory


O/S

LA from
CPU

=<

NO

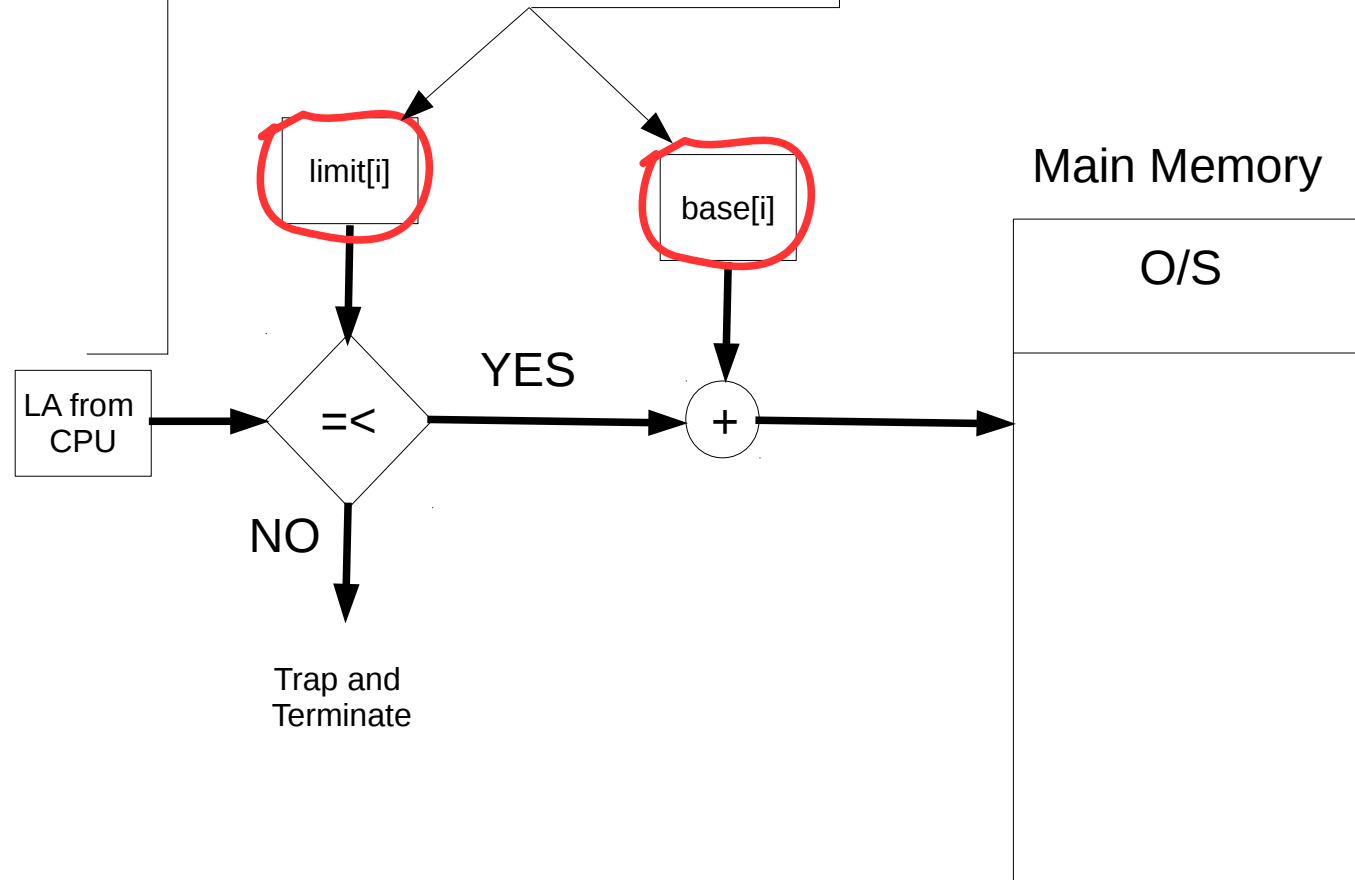
Trap and
Terminate



Process Partition Mapping Table

| BASE | LIMIT |
|------|-------|
| | |
| | |
| | |
| | |
| | |

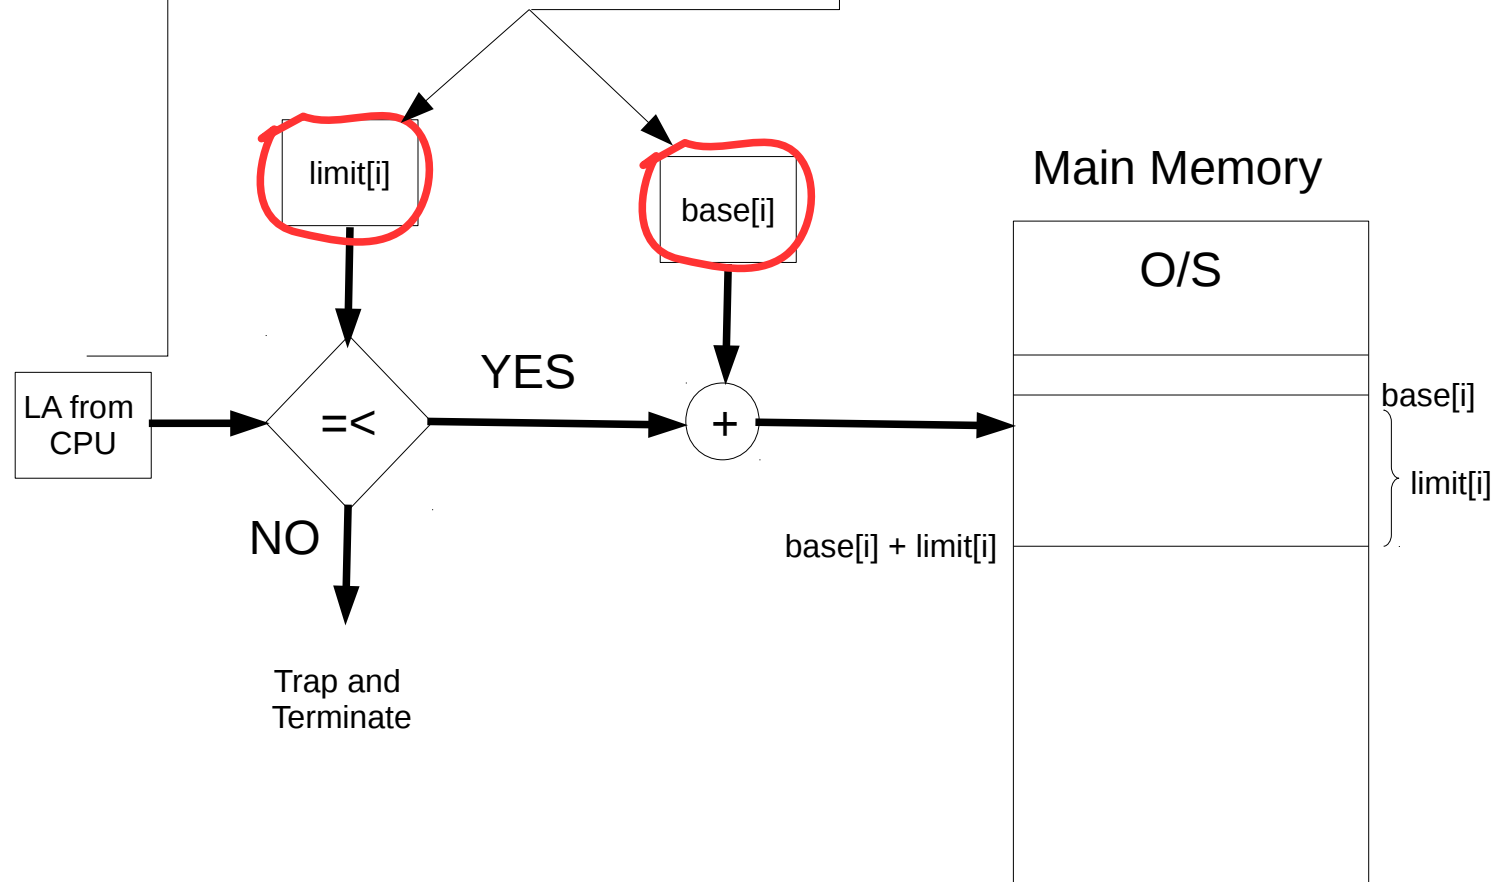
processID



Process Partition Mapping Table

| BASE | LIMIT |
|------|-------|
| | |
| | |
| | |
| | |

processID



Fragmentation in MM

- Memory fragmentation: breaking of memory into non-contiguous blocks that often lead to non-utilisation or wastage of memory
- In this Multi-partition with fixed number processes or fixed-sized partitions: (also called as Multiprogramming with Eixed number of Tasks (MFT))
 - the size of each partition is fixed
 - Partitions could be of different sizes
 - With MFT we have fragmentation
 - Fragmentation: The inability in using memory even though it is empty!
 - Partition gets allocated to a process
 - Several algorithms exist for process-to-partition mapping
 - Different algorithms have different trade-offs particularly w.r.t fragmentation perspective
 - Ofcourse, we need to have an associated data/record
 - Start address
 - Size
 - Allocation status

Partition allocation algorithms

- Two kinds:
 - First-fit algorithm
 - Best-fit algorithm
- First-fit algorithm:
 - Finds the first available partition with the size greater than or equal to the process memory needs
- Best-fit algorithm:
 - Finds the nearest sized partition meeting the process memory size of a process.

Fragmentation types

- Two type
 - Internal memory fragmentation
 - ILLUSTRATE
 - External memory fragmentation
 - ILLUSTRATE
- Best-fit minimises Int. mem. Frag
 - But complexity is high in comparison to FF

Fragmentation in MM

- Memory fragmentation: breaking of memory into non-contiguous blocks that often lead to non-utilisation or wastage of memory
- In this Multi-partition with fixed number of processes:
 - the size of each partition is fixed
 - Partitions could be of different sizes
 - With MFT we have fragmentation
 - Partition gets allocated to a process
 - This record is maintained in Partition Allocation Table

Memory Management

6 November 2018

CS303

Autumn 2018


Memory Mgmt./Multiprogramming With Fixed Number of Tasks (MFT)

- Features of MFT:
 - MM organised into fixed number of partitions
 - Each of fixed size
 - Each partition is defined by lower limit and upper limit
 - LBA/LBR
 - UBA/UBR
 - $\text{Size} = \text{UBA} - \text{LBA}$
 -
 - Each partition gets allocated to a process via algorithms such as
 - First-fit: Less complexity but higher memory fragmentation
 - Best-fit: Increased complexity but decreases fragmentation
- Drawbacks of MFT:
 - As partition number and size are fixed, the max. degree of multiprogramming possible is defined by this number
 - Fragmentation is inevitable since partitions are fixed in advance
 - Fragmentation types: Internal fragmentation and External fragmentation

Multiprog. With Variable Number of Tasks(MVT)

- Features of MVT:
 - Partitions get created of size demanded by processes
 - Having the size as required
 - Each partition is defined by lower limit and upper limit
 - LBA/LBR
 - UBA/UBR
 - $\text{Size} = \text{UBA} - \text{LBA}$
 - Internal fragmentation is eradicated completely, as partitions get created of the required size
- Drawbacks of MFT:
 - External fragmentation is possible


MVT Illustration: Ext. Fragmentation



| ProcessID | SIZE (KB) | CPU Burst |
|-----------|-----------|-----------|
| τ_1 | 100 | 10 |
| τ_2 | 150 | 6 |
| τ_3 | 40 | 25 |
| τ_4 | 90 | 10 |
| τ_5 | 60 | 16 |
| τ_6 | 110 | 10 |



- All processes arriving simultaneously at t_0
- Consider a FCFS

MVT Illustration: Ext. Frag (..cont)



| ProcessID | SIZE (KB) | CPU Burst |
|-----------|-----------|-----------|
| τ_1 | 100 | 10 |
| τ_2 | 150 | 6 |
| τ_3 | 40 | 25 |
| τ_4 | 90 | 10 |
| τ_5 | 60 | 16 |
| τ_6 | 110 | 10 |

MVT Illustration: Ext. Frag (..cont)

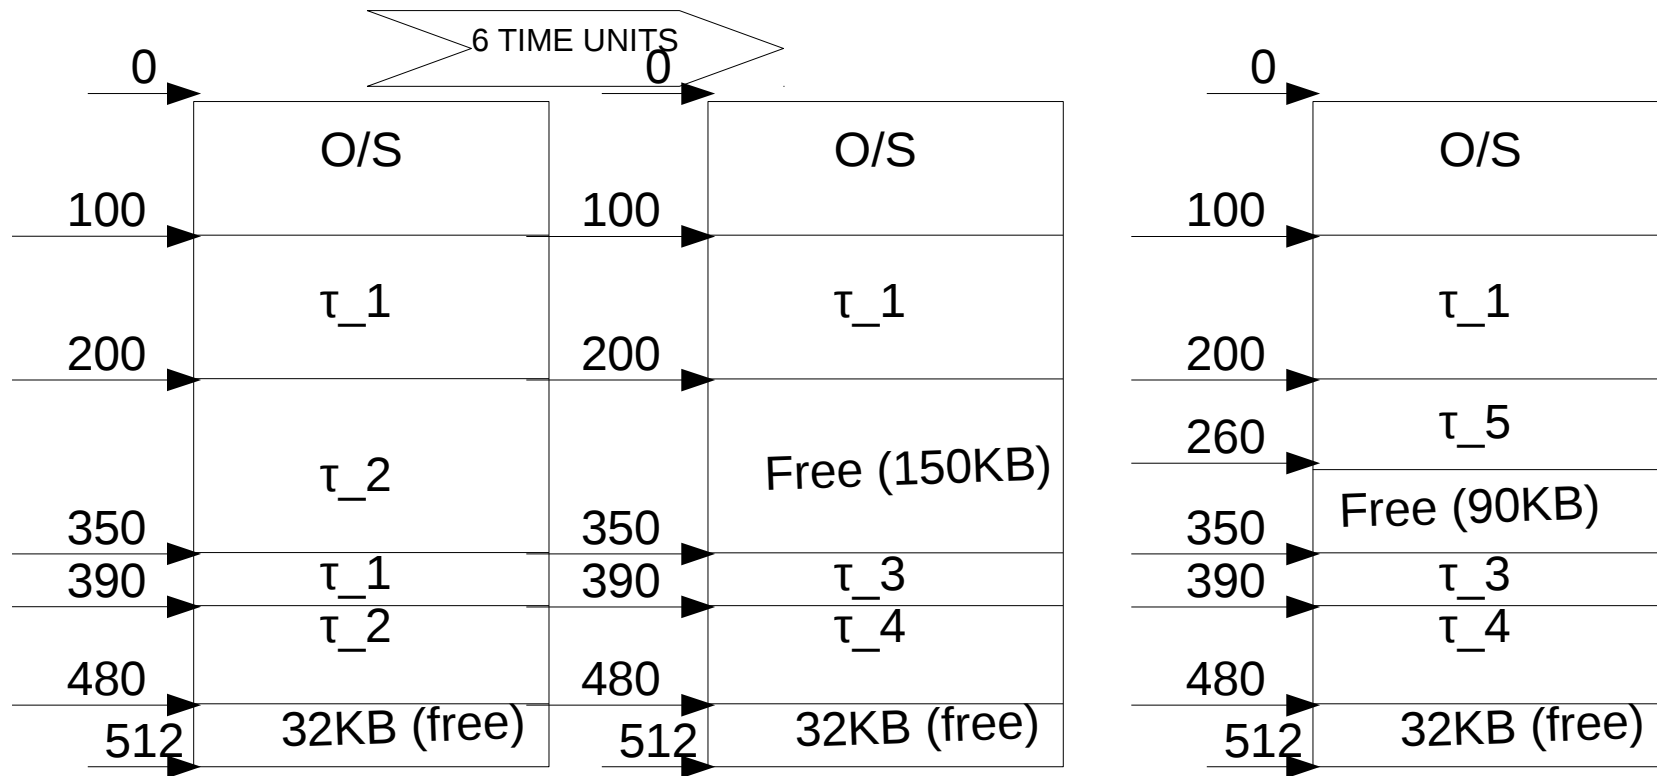


Not useful, as
No process
would fit!!!

| ProcessID | SIZE (KB) | CPU Burst |
|-----------|-----------|-----------|
| τ_1 | 100 | 10 |
| τ_2 | 150 | 6 |
| τ_3 | 40 | 25 |
| τ_4 | 90 | 10 |
| τ_5 | 60 | 16 |
| | 110 | 10 |


EXTERNAL
FRAGMENTATION

MVT Illustration: Ext. Frag (..cont)



| ProcessID | SIZE (KB) | CPU Burst |
|-----------|-----------|-----------|
| τ_1 | 100 | 10 |
| τ_2 | 150 | 6 |
| τ_3 | 40 | 25 |
| τ_4 | 90 | 10 |
| τ_5 | 60 | 16 |
| τ_6 | 110 | 10 |


MVT Illustration: Ext. Frag (..cont)



•

| ProcessID | SIZE (KB) | CPU Burst |
|-----------|-----------|-----------|
| τ_1 | 100 | 10 |
| τ_2 | 150 | 6 |
| τ_3 | 40 | 25 |
| τ_4 | 90 | 10 |
| τ_5 | 60 | 16 |
| τ_6 | 110 | 10 |

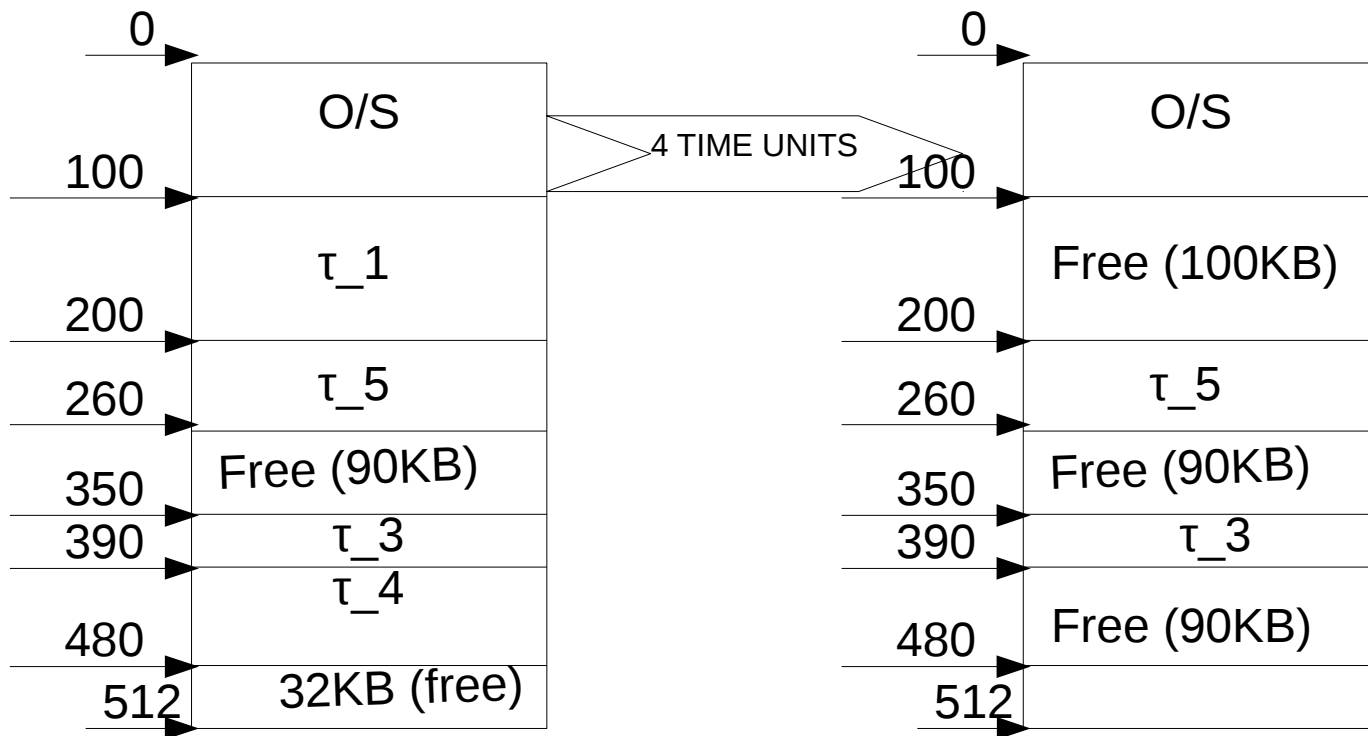
MVT Illustration: Ext. Frag (..cont)



Memory replacement can help?

| ProcessID | SIZE (KB) | CPU Burst |
|-----------|-----------|-----------|
| τ_1 | 100 | 10 |
| τ_2 | 150 | 6 |
| τ_3 | 40 | 25 |
| τ_4 | 90 | 10 |
| τ_5 | 60 | 16 |
| | 110 | 10 |



MVT Illustration: Ext. Frag (..cont)



EXTERNAL
FRAGMENTATION

| ProcessID | SIZE (KB) | CPU Burst |
|-----------|-----------|-----------|
| τ_1 | 100 | 10 |
| τ_2 | 150 | 6 |
| τ_3 | 40 | 25 |
| τ_4 | 90 | 10 |
| τ_5 | 60 | 16 |

MVT Illustration: Ext. Frag (..cont)



Not useful, as
No process
would fit!!!

EXTERNAL
FRAGMENTATION

Though free space is available
As it is not contiguous,
Cannot be allocated to req.
processes!!!!

| ProcessID | SIZE (KB) | CPU Burst |
|-----------|-----------|-----------|
| τ_1 | 100 | 10 |
| τ_2 | 150 | 6 |
| τ_3 | 40 | 25 |
| τ_4 | 90 | 10 |
| τ_5 | 60 | 16 |


Memory Replacement

- If free space is available of required size but is in a discontinuous fashion
 - It cannot be allocated
 - Move all the free areas together, by reallocating the allocated partitions
 - Called **Memory Compaction**
 - Lets move the free areas together to create a larger free area that is contiguous

Memory Compaction

- Memory Compaction algorithm and data str.
 - Keeps track of free areas getting available
 - Keeps track of occupied partitions
 - When advantageous, partition replacement is done
 - A partition table records all this info, which is processed by the algorithm to do the compaction

MVT Illustration: Compaction



Very costly, CPU is completely occupied

| ProcessID | SIZE (KB) | CPU Burst |
|-----------|-----------|-----------|
| τ_1 | 100 | 10 |
| τ_2 | 150 | 6 |
| τ_3 | 40 | 25 |
| τ_4 | 90 | 10 |
| τ_5 | 60 | 16 |
| τ_6 | 110 | 10 |